

02_length_of_stay

July 19, 2021

Nicola De Cristofaro (Matr. 0522500876) Cloud Computing Curriculum

1 Hospital LOS (Length-of-Stay)

First of all what is LOS? **Hospital length-of-stay (LOS)** is defined as the time between hospital admission and discharge measured in days.

1.1 1. Problem Statement

The goal is to create a model that predicts the length-of-stay for each patient at time of admission.

In order to predict hospital LOS, the MIMIC data needed to be separated into terms of: - dependent target variable (length-of-stay in this case) - and independent variables (features) to be used as inputs to the model.

1.2 2. Type of model used for prediction

Since LOS is not a categorical but continuous variable (measured in days), a **regression model** will be used for prediction.

1.3 3. Metrics used for validation

The expected outcome is that the model we use will be better at predicting hospital LOS than the industry standards of **median and average LOS**. The median LOS is simply the median LOS of past admissions to a hospital. Similarly, a second commonly used metric in healthcare is the average, or mean LOS.

So, to measure performance of our model, we'll compare the prediction model against the median and average LOS using the root-mean-square error (RMSE). The RMSE is a commonly used measure of the differences between values predicted by a model and the values observed, where a *lower score implies better accuracy*. For example, a perfect prediction model would have an RMSE of 0.

The RMSE equation for this work is given as follows, where (n) is the number of hospital admission records, (\hat{y}) the prediction LOS, and (y) is the actual LOS.

We could say we have a successful model if its prediction results in a lower RMSE than the average or median models.

There is a multitude of regression models available for predicting LOS. To determine the best regression model between the subset of models that will be evaluated, the **R2 (R-squared)** score will be used.

R Square measures how much variability in dependent variable can be explained by the model. In other words, it is the proportion of the variance in the dependent variable that is predictable from the independent variables. R2 is defined as the following equation where (y_i) is an observed data point, (\hat{y}) is the mean of the observed data, and (f_i) the predicted model value.

Best possible R2 score is 1.0 and a negative value means it is worse than a constant model, average or median in this case.

1.4 4. Features distribution and features engineering

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.linear_model import SGDRegressor
from sklearn.model_selection import GridSearchCV
```

We start importing our baseline dataset extracted selecting only the necessary tables from MIMIC dataset.

```
[4]: # Import baseline dataset constructed in data extraction and preparation phase
admits_patients_diag = pd.read_csv('admits_patients_diag.csv')

#convert dates
admits_patients_diag.admittime = pd.to_datetime(admits_patients_diag.admittime)
admits_patients_diag.disctime = pd.to_datetime(admits_patients_diag.disctime)
admits_patients_diag.deathtime = pd.to_datetime(admits_patients_diag.deathtime)

admits_patients_diag.head()
```

```
[4]: Unnamed: 0  subject_id  hadm_id      admittime      disctime \
0           0    14679932  21038362  2139-09-26  14:16:00  2139-09-28  11:30:00
1           1    15585972  24941086  2123-10-07  23:56:00  2123-10-12  11:22:00
2           2    15078341  23272159  2122-08-28  08:48:00  2122-08-30  12:32:00
3           3    17301855  29732723  2140-06-06  14:23:00  2140-06-08  14:25:00
4           4    17991012  24298836  2181-07-10  20:28:00  2181-07-12  15:49:00

      deathtime admission_type insurance      ethnicity \
0          NaT      ELECTIVE      Other      OTHER/UNKNOWN
```

1	NaT	ELECTIVE	Other		WHITE
2	NaT	ELECTIVE	Other	BLACK/AFRICAN	AMERICAN
3	NaT	ELECTIVE	Other		WHITE
4	NaT	ELECTIVE	Other		WHITE

	died_at_the_hospital	...	injury	mental	misc	muscular	neoplasms	\
0	0	...	2	0	0	0	0	
1	0	...	2	0	0	0	0	
2	0	...	3	0	0	0	0	
3	0	...	2	0	0	0	0	
4	0	...	2	0	0	0	0	

	nervous	pregnancy	prenatal	respiratory	skin
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	1	0	0
4	0	0	0	0	0

[5 rows x 30 columns]

Length of stays computation The LOS is not explicitly expressed as attribute in the admission table, so we have to calculate it. As we said, LOS is defined as the time between admission and discharge from the hospital.

```
[5]: # Create LOS attribute converting timedelta type into float 'days', 86400
      ↪seconds in a day
admits_patients_diag['los'] = (admits_patients_diag['disctime'] -
      ↪admits_patients_diag['admittime']).dt.total_seconds()/86400

# Verify LOS computation
admits_patients_diag[['admittime', 'disctime', 'los']].head()
```

```
[5]:      admittime      disctime      los
0 2139-09-26 14:16:00 2139-09-28 11:30:00 1.884722
1 2123-10-07 23:56:00 2123-10-12 11:22:00 4.476389
2 2122-08-28 08:48:00 2122-08-30 12:32:00 2.155556
3 2140-06-06 14:23:00 2140-06-08 14:25:00 2.001389
4 2181-07-10 20:28:00 2181-07-12 15:49:00 1.806250
```

```
[7]: # We could already have a quick insight on how LOS values are distributed
admits_patients_diag['los'].describe()
```

```
[7]: count      335378.000000
      mean         4.257902
      std         7.223969
      min        -0.945139
```

```

25%          1.129861
50%          2.542361
75%          4.730556
max          1191.416667
Name: los, dtype: float64

```

We noticed that the mean LOS is 4 days, but we noticed also that the min LOS calculated is a negative value, how is it possible that a LOS is negative? Let's see records associated with negative values of LOS:

```
[9]: admits_patients_diag[admits_patients_diag['los'] < 0]
```

```

[9]:      Unnamed: 0  subject_id  hadm_id      admittime \
2359      2359    14556829  27223222  2160-01-04  20:43:00
3416      3416    13362952  25586069  2173-05-02  18:50:00
5927      5927    19649539  20159343  2117-05-27  22:04:00
14137     14137    13604937  22022786  2193-02-16  22:20:00
18846     18846    14316510  27404352  2177-07-26  15:33:00
...      ...      ...      ...      ...
318993    318993    17766453  27023395  2180-07-24  20:51:00
319157    319157    13659453  23252384  2145-05-20  05:08:00
322903    322903    13316652  22658929  2130-05-11  14:55:00
323506    323506    15838787  29390236  2137-07-23  16:58:00
327732    327732    13535122  21247013  2177-08-15  11:56:00

      disctime      deathtime admission_type insurance \
2359  2160-01-04  01:50:00  2160-01-04  01:50:00      ELECTIVE  Medicaid
3416  2173-05-02  09:06:00  2173-05-02  09:06:00      ELECTIVE      Other
5927  2117-05-27  21:16:00                NaT      OBSERVATION      Other
14137  2193-02-16  00:01:00  2193-02-16  00:01:00      EMERGENCY  Medicare
18846  2177-07-26  00:04:00                NaT      OBSERVATION      Other
...      ...      ...      ...      ...
318993  2180-07-24  20:50:00                NaT      OBSERVATION  Medicare
319157  2145-05-20  02:20:00                NaT      EMERGENCY      Other
322903  2130-05-11  02:30:00  2130-05-11  22:23:00      EMERGENCY      Other
323506  2137-07-23  00:01:00  2137-07-23  23:09:00      EMERGENCY      Other
327732  2177-08-15  01:00:00  2177-08-15  01:00:00      EMERGENCY      Other

      ethnicity  died_at_the_hospital  ... mental  misc  \
2359  BLACK/AFRICAN AMERICAN          1  ...      0      0
3416                WHITE            1  ...      0      0
5927                WHITE            0  ...      1      0
14137               WHITE            1  ...      0      0
18846               WHITE            0  ...      0      1
...      ...      ...      ...      ...
318993               WHITE            0  ...      0      1
319157               WHITE            0  ...      1      0
322903               WHITE            1  ...      0      1

```

323506	WHITE	1	...	0	0
327732	ASIAN	1	...	0	2

	muscular	neoplasms	nervous	pregnancy	prenatal	respiratory	skin	\
2359	0	0	0	0	2	0	0	
3416	0	0	0	0	4	0	0	
5927	0	0	0	0	0	0	0	
14137	0	2	0	0	0	0	0	
18846	0	0	0	0	0	0	0	
...	
318993	1	0	0	0	1	0	0	
319157	0	0	0	0	0	0	0	
322903	0	0	4	0	0	0	0	
323506	3	0	3	0	0	1	0	
327732	0	2	0	0	0	2	0	

	los
2359	-0.786806
3416	-0.405556
5927	-0.033333
14137	-0.929861
18846	-0.645139
...	...
318993	-0.000694
319157	-0.116667
322903	-0.517361
323506	-0.706250
327732	-0.455556

[116 rows x 31 columns]

We noticed that rows with negative LOS, usually are related to a time of death before admission, so in this case there is no use to predict LOS, so we drop these rows.

```
[10]: admits_patients_diag = admits_patients_diag[admits_patients_diag['los'] > 0]
      admits_patients_diag.describe()
```

```
[10]:
```

	Unnamed: 0	subject_id	hadm_id	died_at_the_hospital	\
count	335257.000000	3.352570e+05	3.352570e+05	335257.000000	
mean	167692.265262	1.500596e+07	2.500376e+07	0.016719	
std	96815.453052	2.882620e+06	2.888682e+06	0.128215	
min	0.000000	1.000002e+07	2.000002e+07	0.000000	
25%	83845.000000	1.251288e+07	2.250434e+07	0.000000	
50%	167697.000000	1.500984e+07	2.500010e+07	0.000000	
75%	251536.000000	1.749618e+07	2.750696e+07	0.000000	
max	335377.000000	1.999999e+07	2.999983e+07	1.000000	

anchor_age	blood	circulatory	congenital	\
------------	-------	-------------	------------	---

count	335257.000000	335257.000000	335257.000000	335257.000000
mean	50.373036	0.270467	1.455501	0.040459
std	25.542069	0.580229	1.970275	0.239368
min	0.000000	0.000000	0.000000	0.000000
25%	34.000000	0.000000	0.000000	0.000000
50%	54.000000	0.000000	1.000000	0.000000
75%	70.000000	0.000000	2.000000	0.000000
max	91.000000	7.000000	17.000000	11.000000

	digestive	endocrine	...	mental	misc \
count	335257.000000	335257.000000	...	335257.000000	335257.000000
mean	0.586338	1.041777	...	0.597342	0.513469
std	1.037813	1.325580	...	1.027394	0.885857
min	0.000000	0.000000	...	0.000000	0.000000
25%	0.000000	0.000000	...	0.000000	0.000000
50%	0.000000	1.000000	...	0.000000	0.000000
75%	1.000000	2.000000	...	1.000000	1.000000
max	12.000000	12.000000	...	14.000000	13.000000

	muscular	neoplasms	nervous	pregnancy \
count	335257.000000	335257.000000	335257.000000	335257.000000
mean	0.346093	0.221484	0.425918	0.140319
std	0.720684	0.649399	0.795117	0.774798
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	1.000000	0.000000
max	10.000000	11.000000	9.000000	19.000000

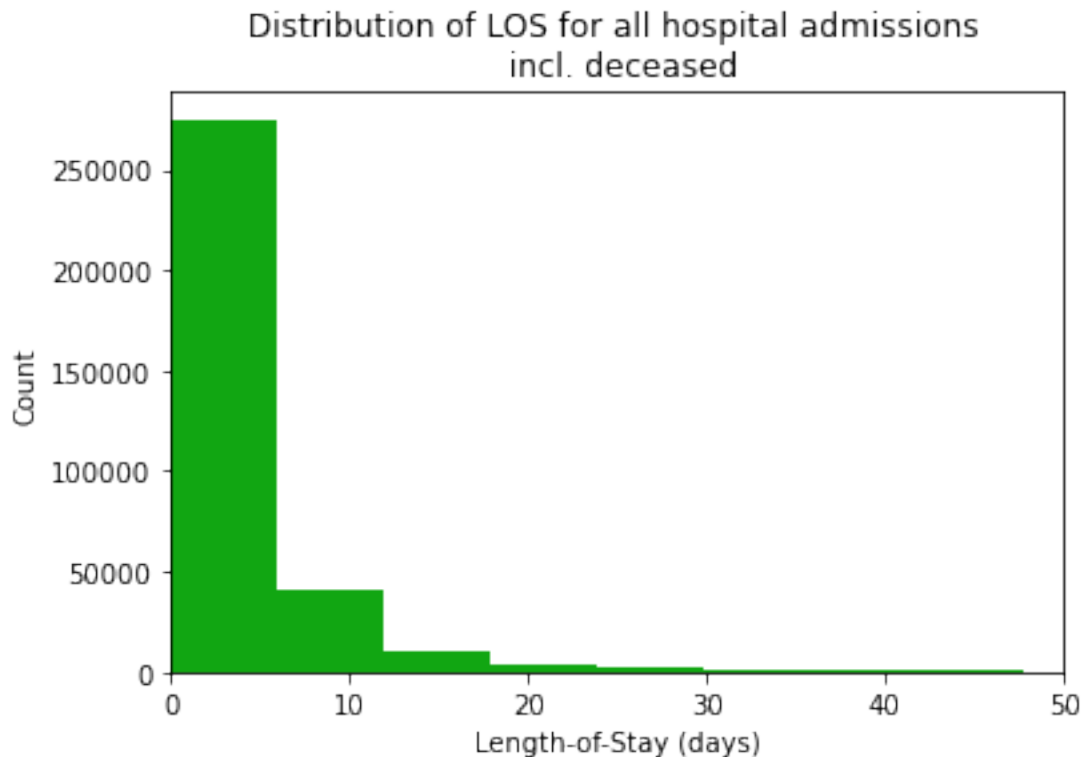
	prenatal	respiratory	skin	los
count	335257.000000	335257.000000	335257.000000	335257.000000
mean	0.237507	0.332038	0.140233	4.259560
std	0.815016	0.705849	0.484663	7.224743
min	0.000000	0.000000	0.000000	0.000694
25%	0.000000	0.000000	0.000000	1.131250
50%	0.000000	0.000000	0.000000	2.543750
75%	0.000000	0.000000	0.000000	4.731944
max	17.000000	9.000000	9.000000	1191.416667

[8 rows x 23 columns]

Now we see how the min value for LOS is not negative anymore. To have a more informative view on the distribution of LOS values we plot those values:

```
[11]: # Plot LOS Distribution
plt.hist(admits_patients_diag['los'], bins=200, color = '#11a612')
plt.xlim(0, 50)
```

```
plt.title('Distribution of LOS for all hospital admissions \n incl. deceased')
plt.ylabel('Count')
plt.xlabel('Length-of-Stay (days)')
plt.tick_params(top=False, right=False)
plt.show();
```



Another thing to consider is admissions of patients who died at the hospital. This kind of admissions resulting in death will be excluded as they would bias the LOS since LOS would be shorter for this group (in data cleaning process this group will be dropped).

```
[12]: print("{} of {} patients died at the hospital".
        ↳format(admits_patients_diag['died_at_the_hospital'].sum(),
        ↳admits_patients_diag['subject_id'].unique()))
```

5605 of 172978 patients died at the hospital

We also said that we'll use the LOS mean and median for comparison and for understand the accuracy of our model. So let's compute these LOS metrics that we'll use later for model evalutaion.

```
[13]: # Hospital LOS metrics for later comparison
actual_mean_los = admits_patients_diag['los'].
        ↳loc[admits_patients_diag['died_at_the_hospital'] == 0].mean()
```

```

actual_median_los = admits_patients_diag['los'].
↳loc[admits_patients_diag['died_at_the_hospital'] == 0].median()

print(actual_mean_los)
print(actual_median_los)

```

```

4.171430522439771
2.5236111111111112

```

```
[ ]: ##### Ethnicity attribute
```

```

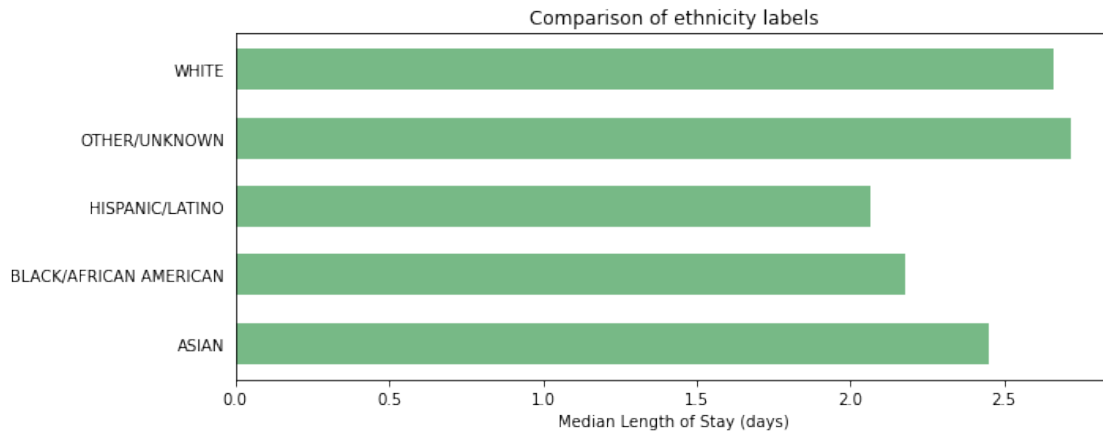
[14]: # Re-usable plotting function
def plot_los_groupby(variable, size=(7,4)):
    '''
    Plot Median LOS by dataframe categorical series name
    '''
    results = admits_patients_diag[[variable, 'los']].groupby(variable).
↳median().reset_index()
    values = list(results['los'].values)
    labels = list(results[variable].values)

    fig, ax = plt.subplots(figsize=size)
    ind = range(len(results))
    ax.barh(ind, values, align='center', height=0.6, color = '#55a868', alpha=0.
↳8)
    ax.set_yticks(ind)
    ax.set_yticklabels(labels)
    ax.set_xlabel('Median Length of Stay (days)')
    ax.tick_params(left=False, top=False, right=False)
    ax.set_title('Comparison of {} labels'.format(variable))

    plt.tight_layout()
    plt.show();

# Look at median LOS for groups ETHNICITY
plot_los_groupby('ethnicity', size=(10,4))

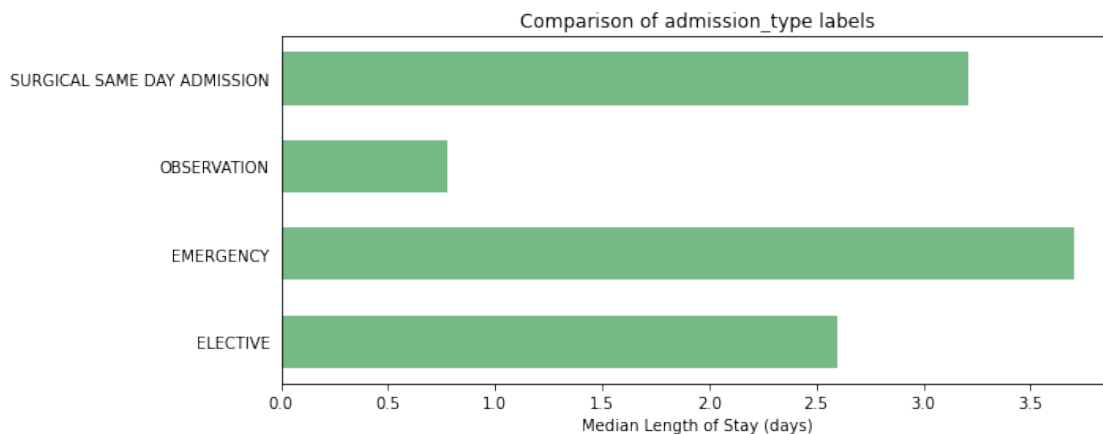
```

To notice that Hispanic/latino patients have the lowest median LOS, even if they are smaller in number in comparison to other ETHNICITY categories.

```
[ ]: ##### ADMISSION_TYPE attribute
```

```
[15]: # Look at median LOS for groups ADMISSION_TYPE
plot_los_groupby('admission_type', size=(10,4))
```



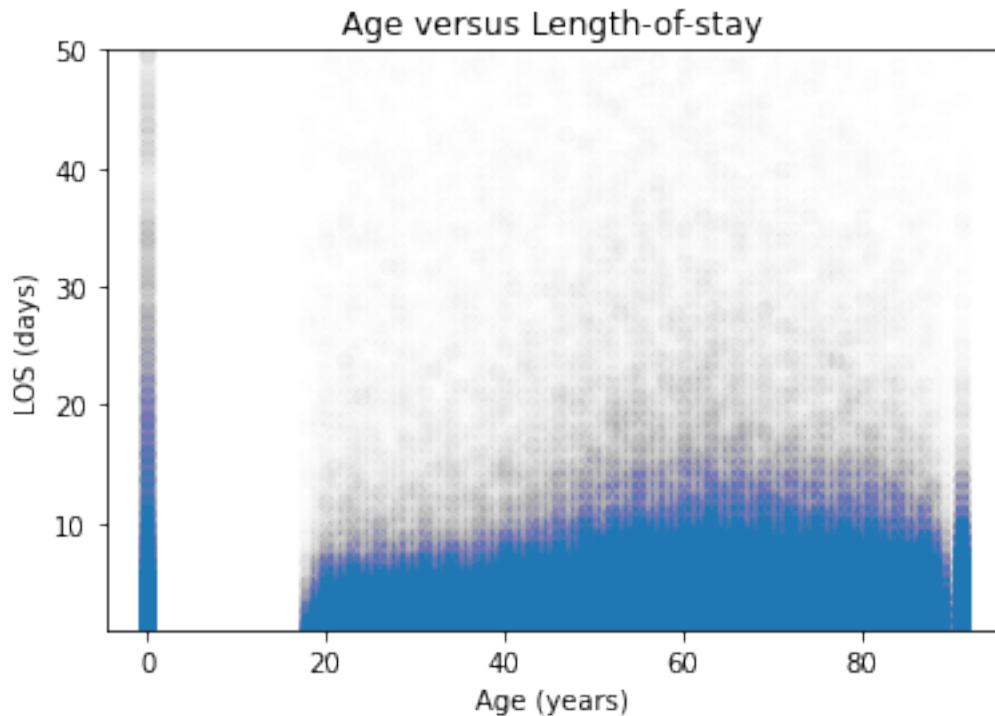
As we could expected *observation* and *elective* admissions have the lowest LOS. This is expected since these are often somewhat planned for and with the risks being understood in comparison to EMERGENCY ADMISSION_TYPE.

AGE attribute Now let's see how the LOS, our current goal, is correlated to ther age of the patients.

```
[16]:
```

```
plt.scatter(admits_patients_diag['anchor_age'], admits_patients_diag['los'],
            alpha=0.005)
plt.ylabel('LOS (days)')
plt.xlabel('Age (years)')
plt.title('Age versus Length-of-stay')
plt.ylim(1, 50)
```

[16]: (1.0, 50.0)



The plot highlights the MIMIC groups of newborns and >89 year olds have higher LOS, and there is an increasing LOS going from 20 toward 80 years old. Because of the discrete-like distribution of data on the extremes of age, it could be useful to convert all ages into the categories of **newborn**, **young adult**, **middle adult**, and **senior** for use in the prediction model.

```
[17]: age_ranges = [(0, 13), (14, 36), (37, 56), (57, 100)]
for num, cat_range in enumerate(age_ranges):
    admits_patients_diag['anchor_age'] = np.
        where(admits_patients_diag['anchor_age'].between(cat_range[0], cat_range[1]),
        num, admits_patients_diag['anchor_age'])

age_dict = {0: 'NEWBORN', 1: 'YOUNG_ADULT', 2: 'MIDDLE_ADULT', 3: 'SENIOR'}
admits_patients_diag['anchor_age'] = admits_patients_diag['anchor_age'].
    replace(age_dict)
```

```
admits_patients_diag.anchor_age.value_counts()
```

```
[17]: SENIOR          155448
      MIDDLE_ADULT    88480
      YOUNG_ADULT     52875
      NEWBORN         38454
      Name: anchor_age, dtype: int64
```

Finally, let's see the distribution of gender in patients in correlation to LOS.

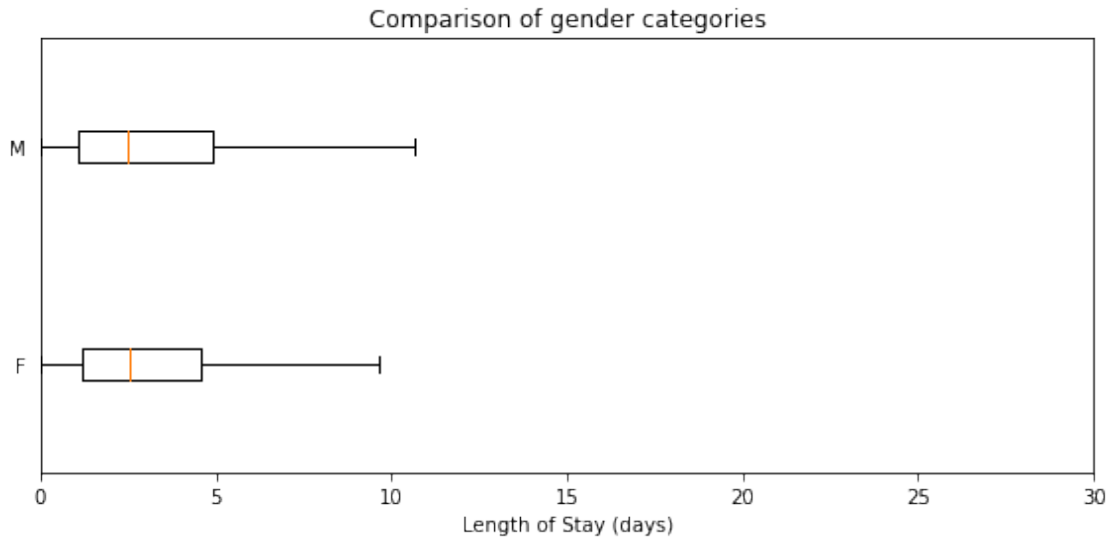
```
[18]: # Re-usable boxplot function
def boxplot_los_groupby(variable, los_range=(-1, 30), size=(8,4)):
    """
    Boxplot of LOS by df categorical series name
    """
    results = admits_patients_diag[[variable, 'los']].groupby(variable).
    ↪median().reset_index()

    categories = results[variable].values.tolist()

    hist_data = []
    for cat in categories:
        hist_data.append(admits_patients_diag['los'].
    ↪loc[admits_patients_diag[variable]==cat].values)

    fig, ax = plt.subplots(figsize=size)
    ax.boxplot(hist_data, 0, '', vert=False)
    ax.set_xlim(los_range)
    ax.set_yticklabels(categories)
    ax.set_xlabel('Length of Stay (days)')
    ax.tick_params(left=False, right=False)
    ax.set_title('Comparison of {} categories'.format(variable))
    plt.tight_layout()
    plt.show();

boxplot_los_groupby('gender', los_range=(0, 30))
```



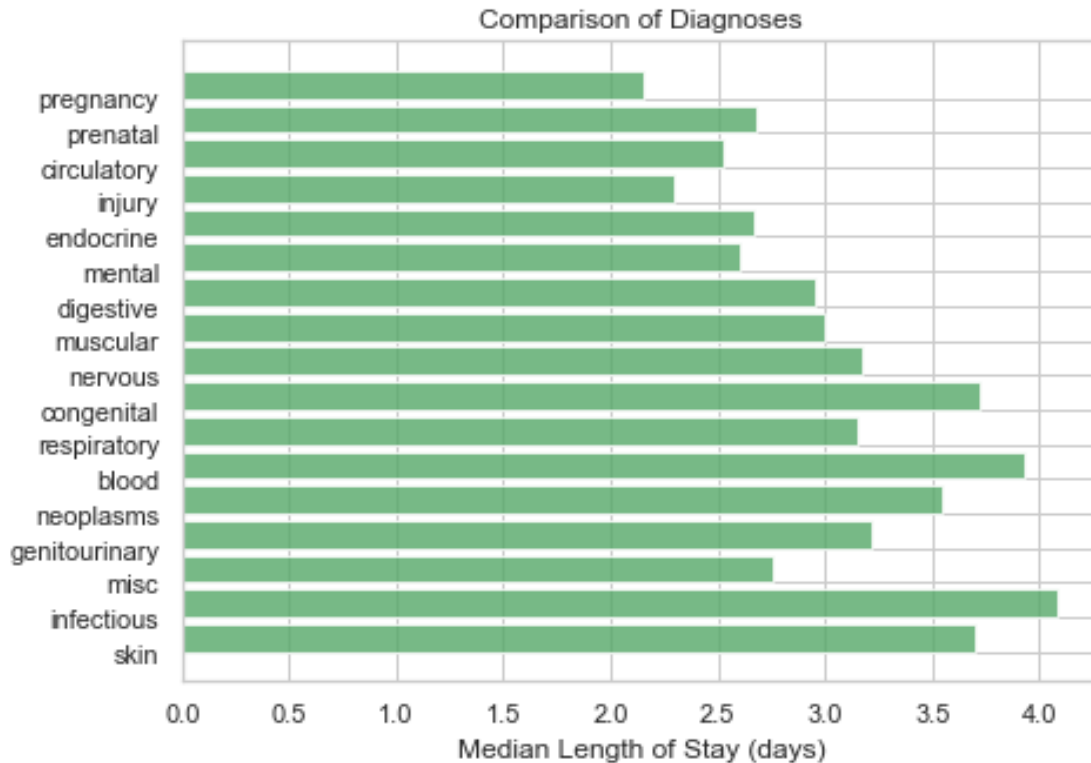
1.5 DIAGNOSIS

[]: Now, let's analyze the diagnosis in correlation to our target LOS.

```
[19]: # Look at the median LOS by diagnosis category
diag_cat_list = ['skin', 'infectious', 'misc', 'genitourinary', 'neoplasms',
↳ 'blood', 'respiratory',
               'congenital', 'nervous', 'muscular', 'digestive', 'mental',
↳ 'endocrine', 'injury',
               'circulatory', 'prenatal', 'pregnancy']

results = []
for variable in diag_cat_list:
    results.append(admits_patients_diag[[variable, 'los']].groupby(variable).
↳ median().reset_index().values[1][1])

sns.set(style="whitegrid")
fig, ax = plt.subplots(figsize=(7,5))
ind = range(len(results))
ax.barh(ind, results, color = '#55a868', alpha=0.8)
ax.set_yticks(ind)
ax.set_yticklabels(diag_cat_list)
ax.set_xlabel('Median Length of Stay (days)')
ax.tick_params(left=False, right=False, top=False)
ax.set_title('Comparison of Diagnoses'.format(variable))
plt.show();
```



Looking at the median LOS for each ICD-9 supercategory shows an important difference between infectious and pregnancy code groups for example. This could already mean that the patients diagnosed an *infection* usually stay longer at the hospital in comparison to other categories.

1.5.1 ICUSTAYS table data extraction

The data in the ICUSTAYS table could be useful because it indicates if a patient during an admission was in an ICU (Intensive Care Unit). This of course could be a factor that could increment the length of stay of patient.

```
[20]: mimic4_path = '../..mimic-iv-1.0/'

# read icustays table
def read_icustays_table(mimic4_path):
    icustays = pd.read_csv(mimic4_path + 'icu/icustays.csv')
    return icustays

icustays = read_icustays_table(mimic4_path)
icustays.head()
```

```
[20]:
```

	subject_id	hadm_id	stay_id	first_careunit	last_careunit	\
0	17867402	24528534	31793211	Trauma SICU (TSICU)	Trauma SICU (TSICU)	
1	14435996	28960964	31983544	Trauma SICU (TSICU)	Trauma SICU (TSICU)	

2	17609946	27385897	33183475	Trauma SICU (TSICU)	Trauma SICU (TSICU)
3	18966770	23483021	34131444	Trauma SICU (TSICU)	Trauma SICU (TSICU)
4	12776735	20817525	34547665	Neuro Stepdown	Neuro Stepdown

	intime		outtime		los
0	2154-03-03	04:11:00	2154-03-04	18:16:56	1.587454
1	2150-06-19	17:57:00	2150-06-22	18:33:54	3.025625
2	2138-02-05	18:54:00	2138-02-15	12:42:05	9.741725
3	2123-10-25	10:35:00	2123-10-25	18:59:47	0.350544
4	2200-07-12	00:33:00	2200-07-13	16:44:40	1.674769

```
[21]: icustays.groupby('first_careunit').median()
```

```
[21]:
```

	subject_id	hadm_id \
first_careunit		
Cardiac Vascular Intensive Care Unit (CVICU)	14892840.0	24981696.0
Coronary Care Unit (CCU)	15027761.0	24967090.5
Medical Intensive Care Unit (MICU)	15024484.0	24970198.0
Medical/Surgical Intensive Care Unit (MICU/SICU)	14995724.0	25014068.0
Neuro Intermediate	14984524.0	25007733.0
Neuro Stepdown	14908467.0	25185450.0
Neuro Surgical Intensive Care Unit (Neuro SICU)	15186511.0	24976128.0
Surgical Intensive Care Unit (SICU)	14971764.0	24869088.0
Trauma SICU (TSICU)	15021710.0	25060781.0

	stay_id	los
first_careunit		
Cardiac Vascular Intensive Care Unit (CVICU)	35007000.0	1.990509
Coronary Care Unit (CCU)	34902339.0	2.011725
Medical Intensive Care Unit (MICU)	35047551.0	1.829190
Medical/Surgical Intensive Care Unit (MICU/SICU)	34964597.0	1.808738
Neuro Intermediate	35070553.0	2.703368
Neuro Stepdown	34947848.0	1.843461
Neuro Surgical Intensive Care Unit (Neuro SICU)	34859288.0	3.646910
Surgical Intensive Care Unit (SICU)	35011708.0	1.974711
Trauma SICU (TSICU)	35047227.0	1.931424

From this statistic we can see how, as far as LOS is concerned, a substantial difference in the median is found only between *Neuro SICU*, *Neuro Intermediate* and the other categories that we can call *Other-ICU*. The *Other-ICU* categories have a very similar median. We can therefore think of simply reducing the categories on three groups: *Neuro SICU*, *Neuro Intermediate* and *Other-ICU* (which includes all the others).

```
[22]:
```

```

icustays['first_careunit'].replace({'Cardiac Vascular Intensive Care Unit_
↳(CVICU)': 'Other-ICU', 'Coronary Care Unit (CCU)': 'Other-ICU', 'Medical_
↳Intensive Care Unit (MICU)': 'Other-ICU', 'Medical/Surgical Intensive Care_
↳Unit (MICU/SICU)': 'Other-ICU', 'Neuro Stepdown': 'Other-ICU', 'Surgical_
↳Intensive Care Unit (SICU)': 'Other-ICU', 'Trauma SICU (TSICU)':_
↳'Other-ICU'}, inplace=True)

icustays['category'] = icustays['first_careunit']
icu_list = icustays.groupby('hadm_id')['category'].apply(list).reset_index()
icu_list.head()

```

```

[22]:      hadm_id      category
0  20000094      [Other-ICU]
1  20000147      [Other-ICU]
2  20000351      [Other-ICU]
3  20000397      [Other-ICU]
4  20000808  [Other-ICU, Other-ICU]

```

```

[23]: icustays['first_careunit'].value_counts()

```

```

[23]: Other-ICU                                72866
      Neuro Surgical Intensive Care Unit (Neuro SICU)    1851
      Neuro Intermediate                                1823
      Name: first_careunit, dtype: int64

```

```

[24]: # Create admission-ICU matrix
icu_item = pd.get_dummies(icu_list['category'].apply(pd.Series).stack()).
↳sum(level=0)
icu_item[icu_item >= 1] = 1
icu_item = icu_item.join(icu_list['hadm_id'], how="outer")
icu_item.head()

```

```

[24]:      Neuro Intermediate      Neuro Surgical Intensive Care Unit (Neuro SICU)  \
0              0              0
1              0              0
2              0              0
3              0              0
4              0              0

      Other-ICU      hadm_id
0              1  20000094
1              1  20000147
2              1  20000351
3              1  20000397
4              1  20000808

```

```
[26]: # Merge ICU data with main dataframe
final_df = admits_patients_diag.merge(icu_item, how='outer', on='hadm_id')
final_df.head()
```

```
[26]: Unnamed: 0  subject_id  hadm_id          admittance          disctime \
0          0.0  14679932.0  21038362  2139-09-26  14:16:00  2139-09-28  11:30:00
1          1.0  15585972.0  24941086  2123-10-07  23:56:00  2123-10-12  11:22:00
2          2.0  15078341.0  23272159  2122-08-28  08:48:00  2122-08-30  12:32:00
3          3.0  17301855.0  29732723  2140-06-06  14:23:00  2140-06-08  14:25:00
4          4.0  17991012.0  24298836  2181-07-10  20:28:00  2181-07-12  15:49:00

      deathtime admission_type insurance          ethnicity \
0          NaT      ELECTIVE      Other      OTHER/UNKNOWN
1          NaT      ELECTIVE      Other              WHITE
2          NaT      ELECTIVE      Other  BLACK/AFRICAN AMERICAN
3          NaT      ELECTIVE      Other              WHITE
4          NaT      ELECTIVE      Other              WHITE

      died_at_the_hospital  ... neoplasms nervous pregnancy prenatal \
0                0.0  ...      0.0      0.0      0.0      0.0
1                0.0  ...      0.0      0.0      0.0      0.0
2                0.0  ...      0.0      0.0      0.0      0.0
3                0.0  ...      0.0      0.0      0.0      1.0
4                0.0  ...      0.0      0.0      0.0      0.0

      respiratory  skin      los  Neuro Intermediate \
0          0.0  0.0  1.884722          NaN
1          0.0  0.0  4.476389          NaN
2          0.0  0.0  2.155556          NaN
3          0.0  0.0  2.001389          NaN
4          0.0  0.0  1.806250          NaN

      Neuro Surgical Intensive Care Unit (Neuro SICU)  Other-ICU
0                NaN                NaN
1                NaN                NaN
2                NaN                NaN
3                NaN                NaN
4                NaN                NaN

[5 rows x 34 columns]
```

```
[27]: final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 362995 entries, 0 to 362994
Data columns (total 34 columns):
#   Column                                     Non-Null Count  Dtype
---  -

```



```

0    Unnamed: 0                                335257 non-null float64
1    subject_id                               335257 non-null float64
2    hadm_id                                   362995 non-null int64
3    admittance                                335257 non-null
datetime64[ns]
4    dischtime                                335257 non-null
datetime64[ns]
5    deathtime                                5605 non-null
datetime64[ns]
6    admission_type                           335257 non-null object
7    insurance                                335257 non-null object
8    ethnicity                                335257 non-null object
9    died_at_the_hospital                     335257 non-null float64
10   gender                                   335257 non-null object
11   anchor_age                               335257 non-null object
12   dod                                       24581 non-null object
13   blood                                    335257 non-null float64
14   circulatory                             335257 non-null float64
15   congenital                              335257 non-null float64
16   digestive                               335257 non-null float64
17   endocrine                               335257 non-null float64
18   genitourinary                           335257 non-null float64
19   infectious                               335257 non-null float64
20   injury                                   335257 non-null float64
21   mental                                   335257 non-null float64
22   misc                                     335257 non-null float64
23   muscular                                335257 non-null float64
24   neoplasms                               335257 non-null float64
25   nervous                                 335257 non-null float64
26   pregnancy                               335257 non-null float64
27   prenatal                                335257 non-null float64
28   respiratory                             335257 non-null float64
29   skin                                    335257 non-null float64
30   los                                     335257 non-null float64
31   Neuro Intermediate                       69211 non-null float64
32   Neuro Surgical Intensive Care Unit (Neuro SICU) 69211 non-null float64
33   Other-ICU                               69211 non-null float64
dtypes: datetime64[ns](3), float64(24), int64(1), object(6)
memory usage: 96.9+ MB

```

```

[30]: # Replace NaNs with 0
final_df['Neuro Intermediate'].fillna(value=0, inplace=True)
final_df['Neuro Surgical Intensive Care Unit (Neuro SICU)'].fillna(value=0,
↪inplace=True)
final_df['Other-ICU'].fillna(value=0, inplace=True)

```

```

[31]: final_df.head()

```

```
[31]: Unnamed: 0  subject_id  hadm_id          admittime          disctime \
0          0.0  14679932.0  21038362  2139-09-26  14:16:00  2139-09-28  11:30:00
1          1.0  15585972.0  24941086  2123-10-07  23:56:00  2123-10-12  11:22:00
2          2.0  15078341.0  23272159  2122-08-28  08:48:00  2122-08-30  12:32:00
3          3.0  17301855.0  29732723  2140-06-06  14:23:00  2140-06-08  14:25:00
4          4.0  17991012.0  24298836  2181-07-10  20:28:00  2181-07-12  15:49:00

      deathtime admission_type insurance          ethnicity \
0          NaT      ELECTIVE      Other      OTHER/UNKNOWN
1          NaT      ELECTIVE      Other              WHITE
2          NaT      ELECTIVE      Other  BLACK/AFRICAN AMERICAN
3          NaT      ELECTIVE      Other              WHITE
4          NaT      ELECTIVE      Other              WHITE

      died_at_the_hospital  ... neoplasms nervous pregnancy prenatal \
0                0.0  ...      0.0      0.0      0.0      0.0
1                0.0  ...      0.0      0.0      0.0      0.0
2                0.0  ...      0.0      0.0      0.0      0.0
3                0.0  ...      0.0      0.0      0.0      1.0
4                0.0  ...      0.0      0.0      0.0      0.0

      respiratory  skin      los  Neuro Intermediate \
0          0.0  0.0  1.884722          0.0
1          0.0  0.0  4.476389          0.0
2          0.0  0.0  2.155556          0.0
3          0.0  0.0  2.001389          0.0
4          0.0  0.0  1.806250          0.0

      Neuro Surgical Intensive Care Unit (Neuro SICU)  Other-ICU
0                0.0      0.0
1                0.0      0.0
2                0.0      0.0
3                0.0      0.0
4                0.0      0.0

[5 rows x 34 columns]
```

1.6 5. Data cleaning

```
[32]: final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 362995 entries, 0 to 362994
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            335257 non-null  float64
1   subject_id                            335257 non-null  float64
```

2	hadm_id	362995 non-null	int64
3	admittime	335257 non-null	
	datetime64[ns]		
4	dischtime	335257 non-null	
	datetime64[ns]		
5	deathtime	5605 non-null	
	datetime64[ns]		
6	admission_type	335257 non-null	object
7	insurance	335257 non-null	object
8	ethnicity	335257 non-null	object
9	died_at_the_hospital	335257 non-null	float64
10	gender	335257 non-null	object
11	anchor_age	335257 non-null	object
12	dod	24581 non-null	object
13	blood	335257 non-null	float64
14	circulatory	335257 non-null	float64
15	congenital	335257 non-null	float64
16	digestive	335257 non-null	float64
17	endocrine	335257 non-null	float64
18	genitourinary	335257 non-null	float64
19	infectious	335257 non-null	float64
20	injury	335257 non-null	float64
21	mental	335257 non-null	float64
22	misc	335257 non-null	float64
23	muscular	335257 non-null	float64
24	neoplasms	335257 non-null	float64
25	nervous	335257 non-null	float64
26	pregnancy	335257 non-null	float64
27	prenatal	335257 non-null	float64
28	respiratory	335257 non-null	float64
29	skin	335257 non-null	float64
30	los	335257 non-null	float64
31	Neuro Intermediate	362995 non-null	float64
32	Neuro Surgical Intensive Care Unit (Neuro SICU)	362995 non-null	float64
33	Other-ICU	362995 non-null	float64

dtypes: datetime64[ns](3), float64(24), int64(1), object(6)

memory usage: 96.9+ MB

```
[34]: # Remove deceased persons as they will skew LOS result
final_df = final_df[final_df['died_at_the_hospital'] == 0.0]
final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 329652 entries, 0 to 335256
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            329652 non-null float64
```

```

1  subject_id          329652 non-null  float64
2  hadm_id            329652 non-null  int64
3  admittime          329652 non-null
datetime64[ns]
4  dischtime          329652 non-null
datetime64[ns]
5  deathtime          0 non-null
datetime64[ns]
6  admission_type     329652 non-null  object
7  insurance          329652 non-null  object
8  ethnicity          329652 non-null  object
9  died_at_the_hospital 329652 non-null  float64
10 gender            329652 non-null  object
11 anchor_age        329652 non-null  object
12 dod              18976 non-null  object
13 blood            329652 non-null  float64
14 circulatory      329652 non-null  float64
15 congenital       329652 non-null  float64
16 digestive        329652 non-null  float64
17 endocrine        329652 non-null  float64
18 genitourinary    329652 non-null  float64
19 infectious       329652 non-null  float64
20 injury           329652 non-null  float64
21 mental           329652 non-null  float64
22 misc             329652 non-null  float64
23 muscular         329652 non-null  float64
24 neoplasms        329652 non-null  float64
25 nervous          329652 non-null  float64
26 pregnancy        329652 non-null  float64
27 prenatal         329652 non-null  float64
28 respiratory      329652 non-null  float64
29 skin             329652 non-null  float64
30 los              329652 non-null  float64
31 Neuro Intermediate 329652 non-null  float64
32 Neuro Surgical Intensive Care Unit (Neuro SICU) 329652 non-null  float64
33 Other-ICU        329652 non-null  float64
dtypes: datetime64[ns](3), float64(24), int64(1), object(6)
memory usage: 88.0+ MB

```

```

[35]: # Remove LOS with negative number
final_df = final_df[final_df['los'] > 0]
final_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 329652 entries, 0 to 335256
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -

```

```

0    Unnamed: 0                                329652 non-null    float64
1    subject_id                               329652 non-null    float64
2    hadm_id                                  329652 non-null    int64
3    admittance                                329652 non-null
datetime64[ns]
4    dischtime                                329652 non-null
datetime64[ns]
5    deathtime                                0 non-null
datetime64[ns]
6    admission_type                           329652 non-null    object
7    insurance                                329652 non-null    object
8    ethnicity                                329652 non-null    object
9    died_at_the_hospital                     329652 non-null    float64
10   gender                                  329652 non-null    object
11   anchor_age                              329652 non-null    object
12   dod                                     18976 non-null     object
13   blood                                   329652 non-null    float64
14   circulatory                             329652 non-null    float64
15   congenital                              329652 non-null    float64
16   digestive                               329652 non-null    float64
17   endocrine                               329652 non-null    float64
18   genitourinary                           329652 non-null    float64
19   infectious                              329652 non-null    float64
20   injury                                  329652 non-null    float64
21   mental                                  329652 non-null    float64
22   misc                                    329652 non-null    float64
23   muscular                                329652 non-null    float64
24   neoplasms                              329652 non-null    float64
25   nervous                                329652 non-null    float64
26   pregnancy                              329652 non-null    float64
27   prenatal                               329652 non-null    float64
28   respiratory                             329652 non-null    float64
29   skin                                    329652 non-null    float64
30   los                                    329652 non-null    float64
31   Neuro Intermediate                      329652 non-null    float64
32   Neuro Surgical Intensive Care Unit (Neuro SICU) 329652 non-null    float64
33   Other-ICU                              329652 non-null    float64
dtypes: datetime64[ns](3), float64(24), int64(1), object(6)
memory usage: 88.0+ MB

```

```

[36]: # Drop unused or no longer needed columns
final_df.drop(columns=['Unnamed: 0', 'subject_id', 'hadm_id', 'admittime',
    → 'dischtime', 'deathtime',
    'died_at_the_hospital', 'dod'], inplace=True)

final_df.info()

```

```
<class 'pandas.core.frame.DataFrame'>
```

Int64Index: 329652 entries, 0 to 335256

Data columns (total 26 columns):

#	Column	Non-Null Count	Dtype
0	admission_type	329652 non-null	object
1	insurance	329652 non-null	object
2	ethnicity	329652 non-null	object
3	gender	329652 non-null	object
4	anchor_age	329652 non-null	object
5	blood	329652 non-null	float64
6	circulatory	329652 non-null	float64
7	congenital	329652 non-null	float64
8	digestive	329652 non-null	float64
9	endocrine	329652 non-null	float64
10	genitourinary	329652 non-null	float64
11	infectious	329652 non-null	float64
12	injury	329652 non-null	float64
13	mental	329652 non-null	float64
14	misc	329652 non-null	float64
15	muscular	329652 non-null	float64
16	neoplasms	329652 non-null	float64
17	nervous	329652 non-null	float64
18	pregnancy	329652 non-null	float64
19	prenatal	329652 non-null	float64
20	respiratory	329652 non-null	float64
21	skin	329652 non-null	float64
22	los	329652 non-null	float64
23	Neuro Intermediate	329652 non-null	float64
24	Neuro Surgical Intensive Care Unit (Neuro SICU)	329652 non-null	float64
25	Other-ICU	329652 non-null	float64

dtypes: float64(21), object(5)

memory usage: 67.9+ MB

```
[37]: final_df.head()
```

```
[37]:  admission_type  insurance      ethnicity  gender  anchor_age  blood  \
0      ELECTIVE    Other      OTHER/UNKNOWN    F      NEWBORN    0.0
1      ELECTIVE    Other              WHITE    F      NEWBORN    0.0
2      ELECTIVE    Other  BLACK/AFRICAN AMERICAN    M      NEWBORN    0.0
3      ELECTIVE    Other              WHITE    F      NEWBORN    0.0
4      ELECTIVE    Other              WHITE    M      NEWBORN    0.0

      circulatory  congenital  digestive  endocrine  ...  neoplasms  nervous  \
0             0.0          1.0         0.0         0.0  ...         0.0        0.0
1             0.0          0.0         0.0         0.0  ...         0.0        0.0
2             0.0          0.0         0.0         0.0  ...         0.0        0.0
3             0.0          0.0         0.0         0.0  ...         0.0        0.0
```

4	0.0	0.0	0.0	0.0	...	0.0	0.0
	pregnancy	prenatal	respiratory	skin	los	Neuro	Intermediate \
0	0.0	0.0	0.0	0.0	1.884722		0.0
1	0.0	0.0	0.0	0.0	4.476389		0.0
2	0.0	0.0	0.0	0.0	2.155556		0.0
3	0.0	1.0	0.0	0.0	2.001389		0.0
4	0.0	0.0	0.0	0.0	1.806250		0.0

	Neuro Surgical Intensive Care Unit (Neuro SICU)	Other-ICU
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

[5 rows x 26 columns]

```
[39]: # Convert gender into numeric boolean attribute
final_df['gender'].replace({'M': 0, 'F': 1}, inplace=True)
final_df.head()
```

```
[39]: admission_type insurance ethnicity gender anchor_age blood \
0 ELECTIVE Other OTHER/UNKNOWN 1 NEWBORN 0.0
1 ELECTIVE Other WHITE 1 NEWBORN 0.0
2 ELECTIVE Other BLACK/AFRICAN AMERICAN 0 NEWBORN 0.0
3 ELECTIVE Other WHITE 1 NEWBORN 0.0
4 ELECTIVE Other WHITE 0 NEWBORN 0.0
```

	circulatory	congenital	digestive	endocrine	...	neoplasms	nervous \
0	0.0	1.0	0.0	0.0	...	0.0	0.0
1	0.0	0.0	0.0	0.0	...	0.0	0.0
2	0.0	0.0	0.0	0.0	...	0.0	0.0
3	0.0	0.0	0.0	0.0	...	0.0	0.0
4	0.0	0.0	0.0	0.0	...	0.0	0.0

	pregnancy	prenatal	respiratory	skin	los	Neuro	Intermediate \
0	0.0	0.0	0.0	0.0	1.884722		0.0
1	0.0	0.0	0.0	0.0	4.476389		0.0
2	0.0	0.0	0.0	0.0	2.155556		0.0
3	0.0	1.0	0.0	0.0	2.001389		0.0
4	0.0	0.0	0.0	0.0	1.806250		0.0

	Neuro Surgical Intensive Care Unit (Neuro SICU)	Other-ICU
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0

3	0.0	0.0
4	0.0	0.0

[5 rows x 26 columns]

```
[40]: # Create dummy columns for categorical variables
prefix_cols = ['ADM', 'INS', 'ETH', 'AGE']
dummy_cols = ['admission_type', 'insurance', 'ethnicity', 'anchor_age']
final_df = pd.get_dummies(final_df, prefix=prefix_cols, columns=dummy_cols)
final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 329652 entries, 0 to 335256
Data columns (total 38 columns):
```

#	Column	Non-Null Count	Dtype
0	gender	329652 non-null	int64
1	blood	329652 non-null	float64
2	circulatory	329652 non-null	float64
3	congenital	329652 non-null	float64
4	digestive	329652 non-null	float64
5	endocrine	329652 non-null	float64
6	genitourinary	329652 non-null	float64
7	infectious	329652 non-null	float64
8	injury	329652 non-null	float64
9	mental	329652 non-null	float64
10	misc	329652 non-null	float64
11	muscular	329652 non-null	float64
12	neoplasms	329652 non-null	float64
13	nervous	329652 non-null	float64
14	pregnancy	329652 non-null	float64
15	prenatal	329652 non-null	float64
16	respiratory	329652 non-null	float64
17	skin	329652 non-null	float64
18	los	329652 non-null	float64
19	Neuro Intermediate	329652 non-null	float64
20	Neuro Surgical Intensive Care Unit (Neuro SICU)	329652 non-null	float64
21	Other-ICU	329652 non-null	float64
22	ADM_ELECTIVE	329652 non-null	uint8
23	ADM_EMERGENCY	329652 non-null	uint8
24	ADM_OBSERVATION	329652 non-null	uint8
25	ADM_SURGICAL SAME DAY ADMISSION	329652 non-null	uint8
26	INS_Medicaid	329652 non-null	uint8
27	INS_Medicare	329652 non-null	uint8
28	INS_Other	329652 non-null	uint8
29	ETH_ASIAN	329652 non-null	uint8
30	ETH_BLACK/AFRICAN AMERICAN	329652 non-null	uint8
31	ETH_HISPANIC/LATINO	329652 non-null	uint8


```

32 ETH_OTHER/UNKNOWN          329652 non-null uint8
33 ETH_WHITE                  329652 non-null uint8
34 AGE_MIDDLE_ADULT           329652 non-null uint8
35 AGE_NEWBORN                 329652 non-null uint8
36 AGE_SENIOR                  329652 non-null uint8
37 AGE_YOUNG_ADULT            329652 non-null uint8
dtypes: float64(21), int64(1), uint8(16)
memory usage: 62.9 MB

```

```
[41]: # Check for any remaining NaNs
final_df.isnull().values.sum()
```

```
[41]: 0
```

The final DataFrame size resulted in 37 feature columns and 1 target column (LOS) with an entry count of 329.652

1.7 6. Prediction Model

We use a **Supervised Learning ML model**. First of all what is it? Supervised learning is defined by its use of labeled datasets to train algorithms that to classify data or predict outcomes accurately. It uses a training set to teach models to yield the desired output. This training dataset includes inputs and correct outputs, which allow the model to learn over time. The algorithm measures its accuracy through the loss function, adjusting until the error has been sufficiently minimized.

Why do we choose it? Because in our case we have the correct output for each dataset entry: LOS (length of stay) and we want to create a model that predicts this output for new entries, in other words that it “generalize well”.

We will implement the supervised learning prediction model using the **Scikit-Learn** machine learning library.

To implement the prediction model, our dataset is splitted into training and test sets at an 80:20 ratio using the scikit-learn *train_test_split* function.

Why split in training and test set? Because to detect a machine learning model behavior, we need to use observations that aren’t used in the training process. Otherwise, the evaluation of the model would be biased as a matter of fact when we build a predictive model, we want the model to work well on data that the model has never seen, so that’s the reason why we use a training set to train the model and a test set to evaluate the model accuracy.

Searching on the Internet for the best train-test ratio, the first answer is 80:20. This means we use 80% of the observations for training and the rest for testing. This approach is taken in this case.

```
[42]: # Target Variable (Length-of-Stay-LOS)
LOS = final_df['los'].values
# Prediction Features
features = final_df.drop(columns=['los'])
```

Using the training set, we’ll fit five different regression models (from the scikit-learn library) using default settings to see what the R2 score comparison looked like.

```
[43]: # Split into training set 80% and test set 20%
X_train, X_test, y_train, y_test = train_test_split(features,
                                                    LOS,
                                                    test_size = .20,
                                                    random_state = 0)

# Show the results of the split
print("Training set has {} samples.".format(X_train.shape[0]))
print("Testing set has {} samples.".format(X_test.shape[0]))
```

Training set has 263721 samples.
Testing set has 65931 samples.

```
[44]: # Regression models used from scikit-learn for comparison
models = [SGDRegressor(random_state = 0),
          GradientBoostingRegressor(random_state = 0),
          LinearRegression(),
          KNeighborsRegressor(),
          RandomForestRegressor(random_state = 0)]

results = {}

for model in models:
    # Instantiate and fit Regressor Model
    reg_model = model
    reg_model.fit(X_train, y_train)

    # Make predictions with model
    y_test_preds = reg_model.predict(X_test)

    # Grab model name and store results associated with model
    name = str(model).split("(")[0]

    results[name] = r2_score(y_test, y_test_preds)
    print('{} done.'.format(name))
```


SGDRegressor done.
GradientBoostingRegressor done.
LinearRegression done.
KNeighborsRegressor done.
RandomForestRegressor done.

```
[45]: # R2 score results
fig, ax = plt.subplots()
ind = range(len(results))
ax.barh(ind, list(results.values()), align='center',
        color = '#55a868', alpha=0.8)
ax.set_yticks(ind)
```

```

ax.set_yticklabels(results.keys())
ax.set_xlabel('R-squared score')
ax.tick_params(left=False, top=False, right=False)
ax.set_title('Comparison of Regression Models')
fig.savefig('images/compare_models.png', bbox_inches = 'tight')

```



The **GradientBoostingRegressor** has the best R2 score of ~48% so we focus on refining this particular model.

```

[46]: # GradientBoostingRegressor will be used as the LOS prediction model
reg_model = GradientBoostingRegressor(random_state=0)
reg_model.fit(X_train, y_train)
y_test_preds = reg_model.predict(X_test)
r2_not_refined = r2_score(y_test, y_test_preds)
print("R2 score is: {:.2f}".format(r2_not_refined))

```

R2 score is: 0.477324

1.8 7. Parameter Tuning

```

[ ]: To refine the GradientBoostingRegressor model, **GridSearchCV** function from
    ↳ scikit-learn is used to test out various permutations of parameters such as
    ↳ *n_estimators, max_depth, and loss*. It helps to loop through predefined
    ↳ hyperparameters and fit your estimator (model) on your training set. So, in
    ↳ the end, we could select the best parameters from the listed hyperparameters.

```

```
[47]: # Split into train 80% and test 20%
X_train, X_test, y_train, y_test = train_test_split(features,
                                                    LOS,
                                                    test_size = .20,
                                                    random_state = 42)

# Set the parameters by cross-validation
#tuned_parameters = [{'n_estimators': [100, 200, 300],
#                        'max_depth' : [2, 3, 4],
#                        'loss': ['ls', 'lad', 'huber']}]
tuned_parameters = [{'n_estimators': [200, 300],
                        'max_depth' : [3, 4],
                        'loss': ['ls', 'lad']}]

# create and fit a ridge regression model, testing each alpha
reg_model = GradientBoostingRegressor()
grid = GridSearchCV(reg_model, tuned_parameters, verbose = 1)
grid.fit(X_train, y_train)
reg_model_optimized = grid.best_estimator_

# summarize the results of the grid search
print(grid.best_score_)
print(grid.best_estimator_)
```

Fitting 5 folds for each of 8 candidates, totalling 40 fits

0.49109741491773634

GradientBoostingRegressor(max_depth=4, n_estimators=300)

Tuned Paramters - *n_estimators*: The number of boosting stages to perform. - *max_depth*: maximum depth of the individual regression estimators. The maximum depth limits the number of nodes in the tree. - *loss*: loss function to be optimized. 'ls' refers to least squares regression. 'lad' (least absolute deviation) is a highly robust loss function solely based on order information of the input variables. 'huber' is a combination of the two.

The best estimator result from GridSearchCV was n_estimators=300, max_depth=4, loss = ls.

```
[48]: y_test_preds = reg_model_optimized.predict(X_test)
r2_optimized = r2_score(y_test, y_test_preds)
print("Optimized R2 score is: {:.2f}".format(r2_optimized))
```

Optimized R2 score is: 0.472332

Parameter tuning didn't improve the R2 score. This could mean that the model is overfitting the training data and can't generalize well on new data. For this reason we continue to use default parameters for GradientBoostingRegressor.

[]: *## 8. Model evaluation and result Discussion*

First of al we could look at what features were most important in predicting ↵
↵hospital length-of-stay when using the gradient boosting regression model.

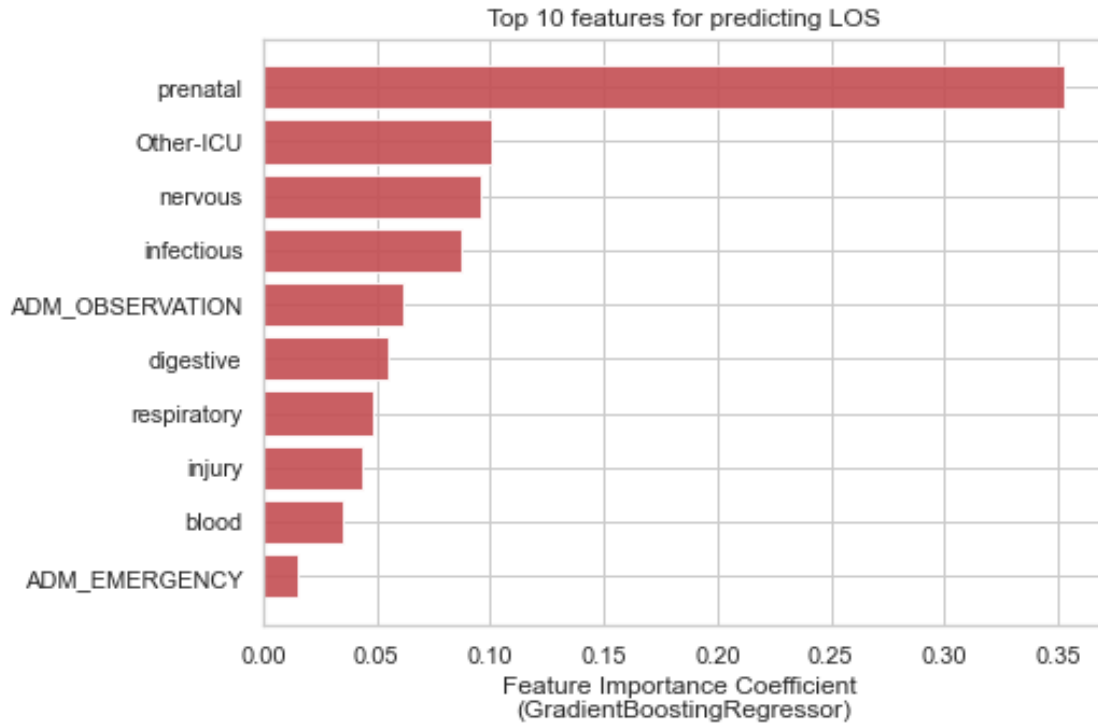
```
[50]: feature_imp = pd.DataFrame(reg_model_optimized.feature_importances_,
                                index = X_train.columns,
                                columns=['importance']).
    ↪sort_values('importance', ascending=False)

feature_imp.head(20)
```

```
[50]:
```

	importance
prenatal	0.352705
Other-ICU	0.100546
nervous	0.095366
infectious	0.087042
ADM_OBSERVATION	0.062016
digestive	0.055106
respiratory	0.048644
injury	0.043253
blood	0.035291
ADM_EMERGENCY	0.015049
neoplasms	0.014466
misc	0.011138
congenital	0.010837
skin	0.010752
ETH_ASIAN	0.008543
circulatory	0.008492
mental	0.007408
endocrine	0.005299
pregnancy	0.004724
ADM_ELECTIVE	0.004527

```
[57]: #Let's plot the top-10 feature importance
feature_imp.index[0:10].tolist()
# Plot feature importance
fig, ax = plt.subplots(figsize=(7, 5))
ind = range(0,10)
ax.barh(ind, feature_imp['importance'].values[0:10],
        align='center', color='#c44e52', alpha=0.9)
ax.set_yticks(ind)
ax.set_yticklabels(feature_imp.index[0:10].tolist())
ax.tick_params(left=False, top=False, right=False)
ax.set_title("Top 10 features for predicting LOS")
ax.set_xlabel('Feature Importance Coefficient \n(GradientBoostingRegressor)')
plt.gca().invert_yaxis()
fig.savefig('images/feature_importance_los_mimic4.png', bbox_inches = 'tight')
```



[]: Diagnoses related to prenatal issues have the highest feature importance coefficient followed by ICU (of general type) admission, nervous and infectious diagnosis. So we could say that, first of all, one of the results is that the *ICD-9 diagnoses categories* are by far the most important features between the features analyzed.

In previous metric section, we said that the RMSE would be used to compare the prediction model versus the industry-standard average and median LOS metrics.

```
[52]: #y_test_preds = reg_model.predict(X_test)

ml_count, md_count, avg_count = 0, 0, 0
ml_days, md_days, avg_days = 0, 0, 0
ml_days_rms, md_days_rms, avg_days_rms = 0, 0, 0

for i in range(y_test_preds.shape[0]):
    ml_model = abs(y_test_preds[i] - y_test[i])
    median_model = abs(actual_median_los - y_test[i])
    average_model = abs(actual_mean_los - y_test[i])

    ml_days += ml_model
    md_days += median_model
    avg_days += average_model
```

```

ml_model_rms = (y_test_preds[i] - y_test[i])**2
median_model_rms = (actual_median_los - y_test[i])**2
average_model_rms = (actual_mean_los - y_test[i])**2

ml_days_rms += ml_model_rms
md_days_rms += median_model_rms
avg_days_rms += average_model_rms

print("Prediction Model days {}".format(ml_days/y_test_preds.shape[0]))
print("Median Model days {}".format(md_days/y_test_preds.shape[0]))
print("Average Model days {}".format(avg_days/y_test_preds.shape[0]))

print("Prediction Model RMS {}".format((ml_days_rms**0.5)/y_test_preds.
    ↳shape[0]))
print("Median Model RMS {}".format((md_days_rms**0.5)/y_test_preds.shape[0]))
print("Average Model RMS {}".format((avg_days_rms**0.5)/y_test_preds.shape[0]))

```

```

Prediction Model days 2.2084883618218436
Median Model days 2.9238206420348725
Average Model days 3.2840947398901026
Prediction Model RMS 0.018297159327123304
Median Model RMS 0.025987491465568332
Average Model RMS 0.025188560774206607

```

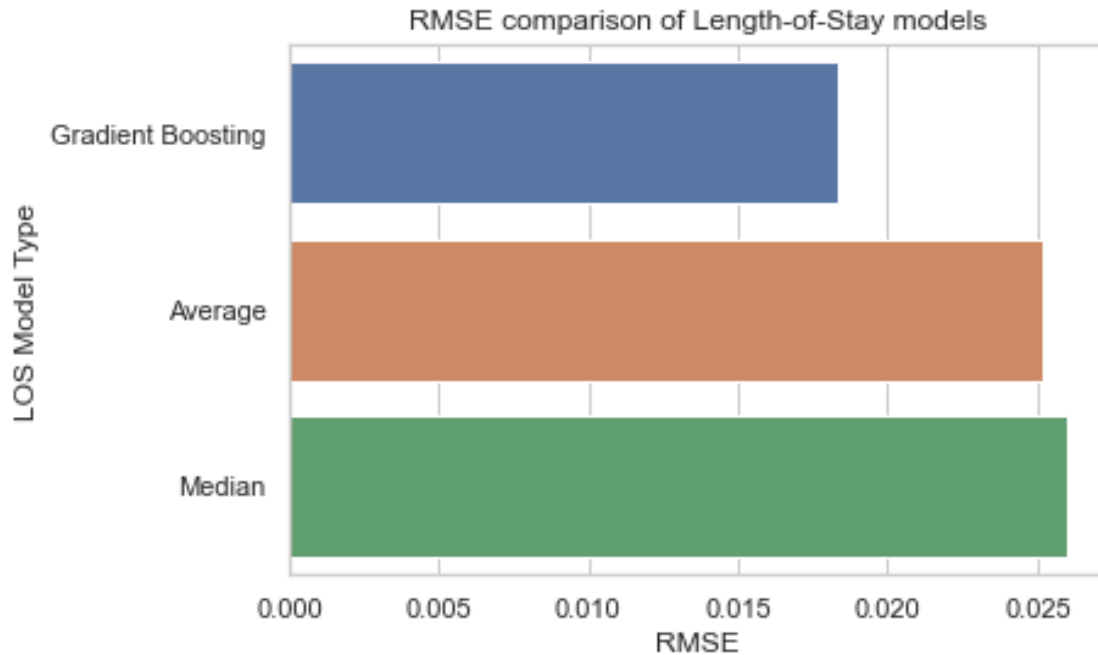
```

[56]: # RMSE plot
data = pd.DataFrame({'RMSE': [(ml_days_rms**0.5)/y_test_preds.shape[0],
                             (avg_days_rms**0.5)/y_test_preds.shape[0],
                             (md_days_rms**0.5)/y_test_preds.shape[0]],
                    'LOS Model Type': ['Gradient Boosting', 'Average', '
    ↳Median'] })

fig, ax = plt.subplots()
ax = sns.barplot(x='RMSE', y='LOS Model Type', data=data)
ax.set_title('RMSE comparison of Length-of-Stay models')
ax.tick_params(top=False, left=False, right=False)

fig.savefig('images/rms_comparison_los_mimic4_01.png', bbox_inches = 'tight')

```



The gradient boosting model RMSE is better even if the percent difference in comparison to the constant average or median models, is not that high (as we can see from the graphic).

Another way to look at the model could be to plot the proportion of accurate predictions in the test set versus an allowed margin of error. Other studies qualify a LOS prediction as correct if it falls within a certain margin of error. Obviously, it follows that as the margin of error allowance increases, so should the proportion of accurate predictions for all models. The gradient boosting prediction model performs better than the other constant models across the margin of error range up to 50%.

```
[55]: # Calculate Proportion of 'accurate' prediction as a function of allowed margin
      ↪ of error
reg_array = []
median_array = []
average_array = []

for i in list(range(6)):
    reg_count, median_count, average_count = 0, 0, 0

    for j in range(y_test_preds.shape[0]):
        # Percent Difference
        reg_model = (y_test_preds[j] - y_test[j])/y_test[j]
        median_model = (actual_median_los - y_test[j])/y_test[j]
        average_model = (actual_mean_los - y_test[j])/y_test[j]
        if abs(reg_model) < i/10:
            reg_count += 1
```



```

    if abs(median_model) < i/10:
        median_count += 1
    if abs(average_model) < i/10:
        average_count += 1

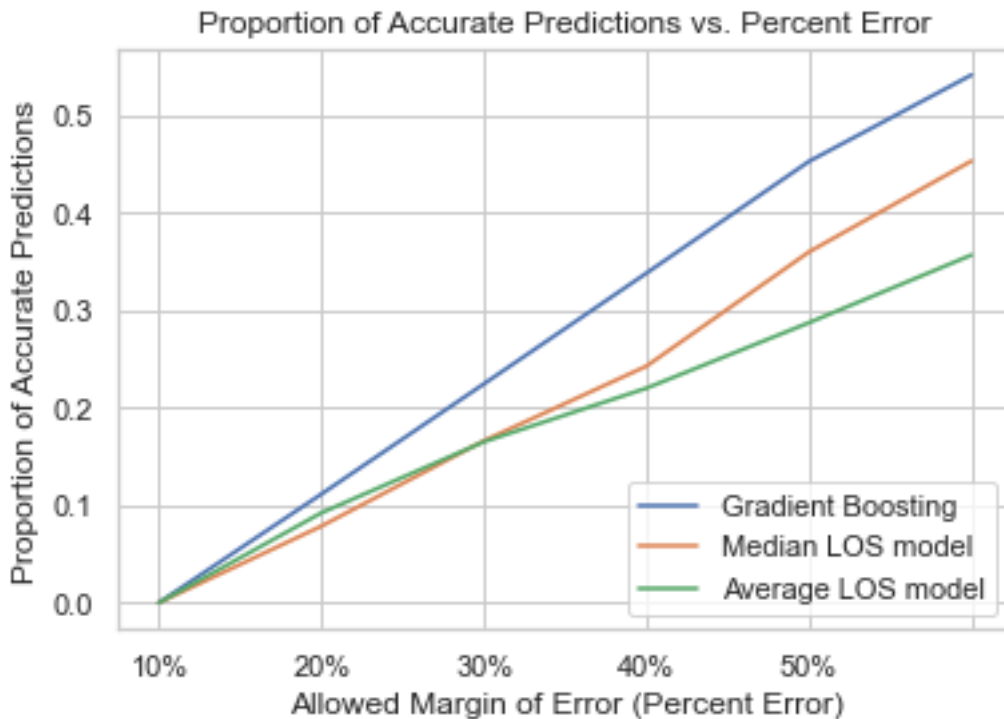
    reg_array.append((reg_count/y_test_preds.shape[0]))
    median_array.append((median_count/y_test_preds.shape[0]))
    average_array.append((average_count/y_test_preds.shape[0]))

# Plot proportion of 'accurate' prediction as a function of allowed margin of
↪ error
fig, ax = plt.subplots()
ax.plot(reg_array, label='Gradient Boosting')
ax.plot(median_array, label='Median LOS model')
ax.plot(average_array, label='Average LOS model')
ax.set_title('Proportion of Accurate Predictions vs. Percent Error')
ax.set_xlabel('Allowed Margin of Error (Percent Error)')
ax.set_ylabel('Proportion of Accurate Predictions')
ax.set_xticklabels(['0%', '10%', '20%', '30%', '40%', '50%'])
ax.legend(loc='lower right');
ax.tick_params(top=False, right=False)
fig.savefig('images/rms_comparison_los_mimic4_02.png', bbox_inches = 'tight')

```

<ipython-input-55-86a6c24f915a>:33: UserWarning: FixedFormatter should only be used together with FixedLocator

```
ax.set_xticklabels(['0%', '10%', '20%', '30%', '40%', '50%'])
```



1.9 Conclusions for LOS (Length-of-stay)

Hospital stays cost the health system at least a big amount of money. U.S. Hospital for example spends \$377.5 billion per year in the health system and recent Medicare legislation standardizes payments for procedures performed, regardless of the number of days a patient spends in the hospital.

This incentivizes hospitals to identify patients of high LOS risk at the time of admission. Once identified, patients with high LOS risk can have their treatment plan optimized to minimize LOS and lower the chance of getting a hospital-acquired condition. Another benefit is that prior knowledge of LOS can aid in logistics such as room and bed allocation planning.