

03_in_hospital_mortality

July 19, 2021

Nicola De Cristofaro (Matr. 0522500876) Cloud Computing Curriculum

1 In-Hospital Mortality

First of all what do we mean by “In-Hospital Mortality”? In-Hospital Mortality refers to a death occurred during the hospital stay.

1.1 1. Problem Statement

The goal is to create a model that predicts if a patient will incur in death during its hospital stay.

1.2 2. Type of model used for prediction

“In-Hospital Mortality” is a categorical attribute (YES,NO), so normally a classification model has to be used, but in this case we have a binary classifier so we could also treat the two possible output categories as numeric (YES=1, NO=0) and use a **regression model**. We will do Regression in our task.

1.3 3. Metrics used for validation

So, to measure performance of our model, we'll compute the root-mean-square error (RMSE). The RMSE is a commonly used measure of the differences between values predicted by a model and the values observed, where a *lower score implies better accuracy*. For example, a perfect prediction model would have an RMSE of 0.

The RMSE equation for this work is given as follows, where (n) is the number of hospital admission records, (\hat{y}) the prediction In-Hospital Mortality, and (y) is the actual In-Hospital Mortality.

To determine the best regression model between the subset of models that will be evaluated, the **R2 (R-squared)** score will be used.

R Square measures how much variability in dependent variable can be explained by the model. In other words, it is the proportion of the variance in the dependent variable that is predictable from the independent variables. R2 is defined as the following equation where (y_i) is an observed data point, (\hat{y}) is the mean of the observed data, and (f_i) the predicted model value.

Best possible R2 score is 1.0.

```
[136]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

We start importing our baseline dataset extracted selecting only the necessary tables from MIMIC dataset.

```
[137]: # Import baseline dataset constructed in data extraction and preparation phase
admits_patients_diag = pd.read_csv('admits_patients_diag.csv')

#convert dates
admits_patients_diag.admittime = pd.to_datetime(admits_patients_diag.admittime)
admits_patients_diag.disctime = pd.to_datetime(admits_patients_diag.disctime)
admits_patients_diag.deathtime = pd.to_datetime(admits_patients_diag.deathtime)

admits_patients_diag.head()
```

```
[137]: Unnamed: 0  subject_id  hadm_id      admittime      disctime \
0          0    14679932  21038362  2139-09-26  14:16:00  2139-09-28  11:30:00
1          1    15585972  24941086  2123-10-07  23:56:00  2123-10-12  11:22:00
2          2    15078341  23272159  2122-08-28  08:48:00  2122-08-30  12:32:00
3          3    17301855  29732723  2140-06-06  14:23:00  2140-06-08  14:25:00
4          4    17991012  24298836  2181-07-10  20:28:00  2181-07-12  15:49:00

      deathtime admission_type insurance      ethnicity \
0          NaT      ELECTIVE      Other      OTHER/UNKNOWN
1          NaT      ELECTIVE      Other              WHITE
2          NaT      ELECTIVE      Other  BLACK/AFRICAN AMERICAN
3          NaT      ELECTIVE      Other              WHITE
4          NaT      ELECTIVE      Other              WHITE

      died_at_the_hospital  ... injury  mental misc  muscular  neoplasms \
0                0  ...      2      0      0          0          0
1                0  ...      2      0      0          0          0
2                0  ...      3      0      0          0          0
3                0  ...      2      0      0          0          0
4                0  ...      2      0      0          0          0

      nervous  pregnancy  prenatal  respiratory  skin
0           0           0          0            0      0
1           0           0          0            0      0
2           0           0          0            0      0
3           0           0          1            0      0
4           0           0          0            0      0

[5 rows x 30 columns]
```

```
[138]: admits_patients_diag.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 335378 entries, 0 to 335377
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            335378 non-null  int64
1   subject_id            335378 non-null  int64
2   hadm_id               335378 non-null  int64
3   admittime             335378 non-null  datetime64[ns]
4   disctime              335378 non-null  datetime64[ns]
5   deathtime             5670 non-null    datetime64[ns]
6   admission_type        335378 non-null  object
7   insurance             335378 non-null  object
8   ethnicity             335378 non-null  object
9   died_at_the_hospital  335378 non-null  int64
10  gender                335378 non-null  object
11  anchor_age            335378 non-null  int64
12  dod                   24651 non-null   object
13  blood                 335378 non-null  int64
14  circulatory           335378 non-null  int64
15  congenital            335378 non-null  int64
16  digestive              335378 non-null  int64
17  endocrine              335378 non-null  int64
18  genitourinary         335378 non-null  int64
19  infectious             335378 non-null  int64
20  injury                335378 non-null  int64
21  mental                335378 non-null  int64
22  misc                  335378 non-null  int64
23  muscular              335378 non-null  int64
24  neoplasms             335378 non-null  int64
25  nervous               335378 non-null  int64
26  pregnancy             335378 non-null  int64
27  prenatal              335378 non-null  int64
28  respiratory           335378 non-null  int64
29  skin                  335378 non-null  int64
dtypes: datetime64[ns](3), int64(22), object(5)
memory usage: 76.8+ MB

```

```

[139]: # Create LOS attribute converting timedelta type into float 'days', 86400
        ↪seconds in a day
admits_patients_diag['los'] = (admits_patients_diag['disctime'] -
        ↪admits_patients_diag['admittime']).dt.total_seconds()/86400

# Verify LOS computation
admits_patients_diag[['admittime', 'disctime', 'los']].head()

```

```
[139]:      admittime      disctime      los
0 2139-09-26 14:16:00 2139-09-28 11:30:00 1.884722
1 2123-10-07 23:56:00 2123-10-12 11:22:00 4.476389
2 2122-08-28 08:48:00 2122-08-30 12:32:00 2.155556
3 2140-06-06 14:23:00 2140-06-08 14:25:00 2.001389
4 2181-07-10 20:28:00 2181-07-12 15:49:00 1.806250
```

We have computed the length of stay for each admission because there are some LOS less than 0 (negative), they could refer to a patient's death before the admission; this kind of entry has to be dropped.

```
[140]: # Remove LOS with negative number
admits_patients_diag = admits_patients_diag[admits_patients_diag['los'] > 0]
admits_patients_diag.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 335257 entries, 0 to 335377
Data columns (total 31 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Unnamed: 0          335257 non-null  int64
1   subject_id          335257 non-null  int64
2   hadm_id             335257 non-null  int64
3   admittime           335257 non-null  datetime64[ns]
4   disctime            335257 non-null  datetime64[ns]
5   deathtime           5605 non-null    datetime64[ns]
6   admission_type      335257 non-null  object
7   insurance            335257 non-null  object
8   ethnicity            335257 non-null  object
9   died_at_the_hospital 335257 non-null  int64
10  gender               335257 non-null  object
11  anchor_age           335257 non-null  int64
12  dod                  24581 non-null   object
13  blood                335257 non-null  int64
14  circulatory          335257 non-null  int64
15  congenital           335257 non-null  int64
16  digestive            335257 non-null  int64
17  endocrine            335257 non-null  int64
18  genitourinary        335257 non-null  int64
19  infectious           335257 non-null  int64
20  injury               335257 non-null  int64
21  mental               335257 non-null  int64
22  misc                 335257 non-null  int64
23  muscular             335257 non-null  int64
24  neoplasms            335257 non-null  int64
25  nervous              335257 non-null  int64
26  pregnancy            335257 non-null  int64
27  prenatal             335257 non-null  int64
```

```

28 respiratory          335257 non-null  int64
29 skin                 335257 non-null  int64
30 los                  335257 non-null  float64
dtypes: datetime64[ns](3), float64(1), int64(22), object(5)
memory usage: 81.8+ MB

```

```
[141]: admits_patients_diag.died_at_the_hospital.value_counts()
```

```

[141]: 0    329652
      1     5605
      Name: died_at_the_hospital, dtype: int64

```

```

[142]: # Plot LOS Distribution
plt.hist(admits_patients_diag['died_at_the_hospital'], bins=3, color = '#11a612')
plt.title('Distribution of Died_at_th_Hospital for all hospital admissions')
plt.ylabel('Count')
plt.xlabel('Died_at_the_Hospital')
plt.xticks(range(0,2))
plt.show();

```



Died_at_the_Hospital = 1 means that the patients is died during the hospital stay. As we can see, luckily we have a small number of death compared to survivors.

```
[143]: print('Number of positive samples:', (admits_patients_diag.died_at_the_hospital_
↪== 1).sum()) # died at hospital
print('Number of negative samples:', (admits_patients_diag.
↪died_at_the_hospital == 0).sum()) # not died at hospital
print('Total samples:', len(admits_patients_diag))
```

Number of positive samples: 5605

Number of negative samples: 329652

Total samples: 335257

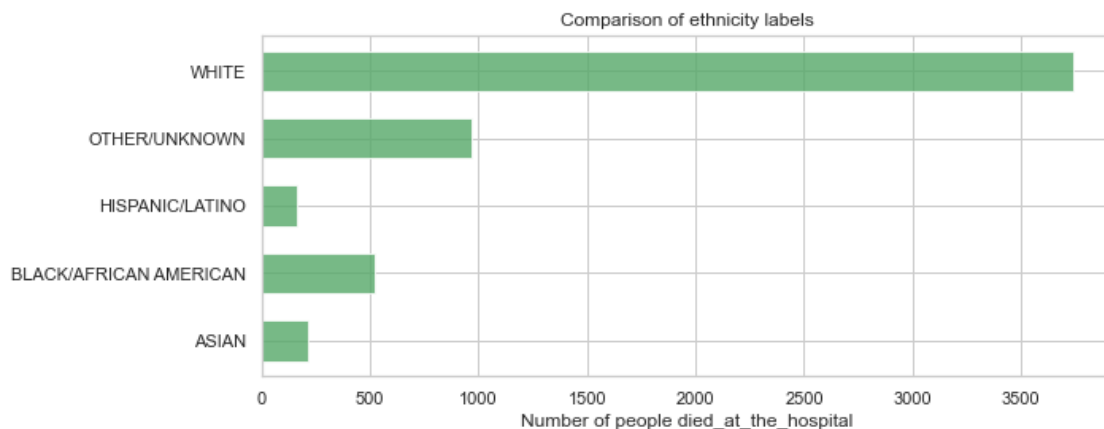
1.3.1 Died_at_the_Hospital VS Ethnicity

```
[144]: # Re-usable plotting function
def plot_los_groupby(variable, size=(7,4)):
    results = admits_patients_diag[[variable, 'died_at_the_hospital']]
    ↪groupby(variable).sum().reset_index()
    values = list(results['died_at_the_hospital'].values)
    labels = list(results[variable].values)

    fig, ax = plt.subplots(figsize=size)
    ind = range(len(results))
    ax.barh(ind, values, align='center', height=0.6, color = '#55a868', alpha=0.
    ↪8)

    ax.set_yticks(ind)
    ax.set_yticklabels(labels)
    ax.set_xlabel('Number of people died_at_the_hospital')
    ax.tick_params(left=False, top=False, right=False)
    ax.set_title('Comparison of {} labels'.format(variable))
    plt.tight_layout()
    plt.show();

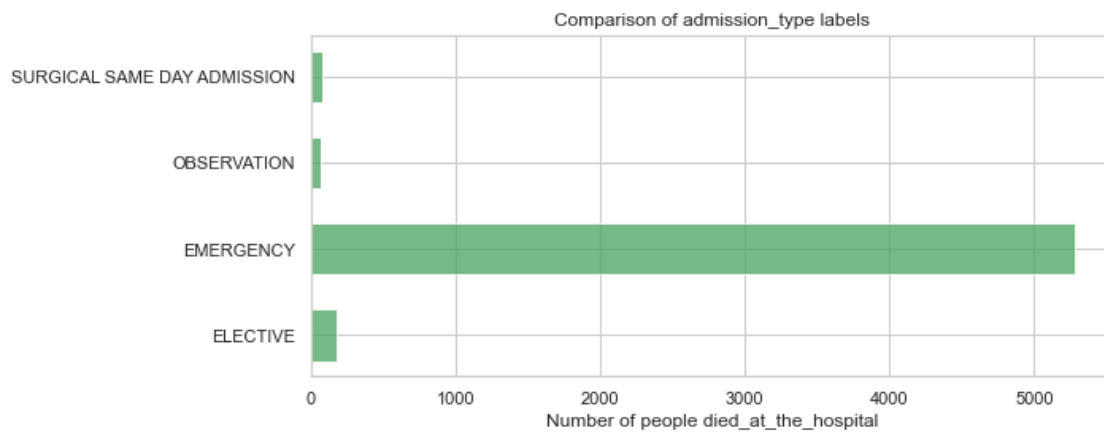
# Look at median LOS for groups ETHNICITY
plot_los_groupby('ethnicity', size=(10,4))
```



We can notice that Asian and Hispanic/latino people usually do not incur in death during their hospital stays.

1.3.2 Died_at_the_Hospital VS Admission_Type


```
[145]: plot_los_groupby('admission_type', size=(10,4))
```



As we expect the biggest number of death is for patients admitted with emergency.

1.3.3 Died_at_the_Hospital VS Age

```
[146]: plot_los_groupby('anchor_age', size=(10,10))
```



Because of the discrete-like distribution of data on the extremes of age (0 and >89), it could be useful to convert all ages into the categories of **newborn**, **young adult**, **middle adult**, and **senior** for use in the prediction model.

```
[147]: age_ranges = [(0, 13), (14, 36), (37, 56), (57, 100)]
for num, cat_range in enumerate(age_ranges):
    admits_patients_diag['anchor_age'] = np.
    ↳where(admits_patients_diag['anchor_age'].between(cat_range[0],cat_range[1]),
    ↳num, admits_patients_diag['anchor_age'])

age_dict = {0: 'NEWBORN', 1: 'YOUNG_ADULT', 2: 'MIDDLE_ADULT', 3: 'SENIOR'}
admits_patients_diag['anchor_age'] = admits_patients_diag['anchor_age'].
    ↳replace(age_dict)
```



```
admits_patients_diag.anchor_age.value_counts()
```

```
[147]: SENIOR          155448
      MIDDLE_ADULT    88480
      YOUNG_ADULT     52875
      NEWBORN         38454
      Name: anchor_age, dtype: int64
```


Now, let's analyze the diagnosis in correlation to our target LOS.

```
[148]: import seaborn as sns

# Look at the median LOS by diagnosis category
diag_cat_list = ['skin', 'infectious', 'misc', 'genitourinary', 'neoplasms', '
↳ 'blood', 'respiratory',
                  'congenital', 'nervous', 'muscular', 'digestive', 'mental', '
↳ 'endocrine', 'injury',
                  'circulatory', 'prenatal', 'pregnancy']

results = []
for variable in diag_cat_list:
    results.append(admits_patients_diag[[variable, 'died_at_the_hospital']].
↳ groupby(variable).sum().reset_index().values[1][1])

sns.set(style="whitegrid")
fig, ax = plt.subplots(figsize=(7,5))
ind = range(len(results))
ax.barh(ind, results, color = '#55a868', alpha=0.8)
ax.set_yticks(ind)
ax.set_yticklabels(diag_cat_list)
ax.set_xlabel('Sum of people Died_at_the_Hospital')
ax.tick_params(left=False, right=False, top=False)
ax.set_title('Comparison of Diagnoses'.format(variable))
plt.show();
```



We can have interesting insights from this plot: - the biggest number of deaths in hospital is for people diagnosed in the category **genitourinary** - the lowest number of deaths in hospital (0 in the subset considered) is for people in **pregnancy**

1.3.4 ICUSTAYS table data extraction

The data in the ICUSTAYS table could be useful because indicates if a patient during an admission was in an ICU (Intensive Care Unit). This of course could be a factor that could determine the death of the patient at the hospital.

```
[149]: mimic4_path = '../mimic-iv-1.0/'

# read icustays table
def read_icustays_table(mimic4_path):
    icustays = pd.read_csv(mimic4_path + 'icu/icustays.csv')
    return icustays

icustays = read_icustays_table(mimic4_path)
icustays.head()
```

	subject_id	hadm_id	stay_id	first_careunit	last_careunit \
0	17867402	24528534	31793211	Trauma SICU (TSICU)	Trauma SICU (TSICU)
1	14435996	28960964	31983544	Trauma SICU (TSICU)	Trauma SICU (TSICU)

2	17609946	27385897	33183475	Trauma SICU (TSICU)	Trauma SICU (TSICU)
3	18966770	23483021	34131444	Trauma SICU (TSICU)	Trauma SICU (TSICU)
4	12776735	20817525	34547665	Neuro Stepdown	Neuro Stepdown

		intime		outtime	los
0	2154-03-03	04:11:00	2154-03-04	18:16:56	1.587454
1	2150-06-19	17:57:00	2150-06-22	18:33:54	3.025625
2	2138-02-05	18:54:00	2138-02-15	12:42:05	9.741725
3	2123-10-25	10:35:00	2123-10-25	18:59:47	0.350544
4	2200-07-12	00:33:00	2200-07-13	16:44:40	1.674769

```
[150]: icustays['category'] = icustays['first_careunit']
icu_list = icustays.groupby('hadm_id')['category'].apply(list).reset_index()
icu_list.head()
```

```
[150]:      hadm_id      category
0  20000094  [Coronary Care Unit (CCU)]
1  20000147  [Coronary Care Unit (CCU)]
2  20000351  [Medical/Surgical Intensive Care Unit (MICU/SI...
3  20000397  [Coronary Care Unit (CCU)]
4  20000808  [Surgical Intensive Care Unit (SICU), Surgical...
```

```
[151]: icustays['first_careunit'].value_counts()
```

```
[151]: Medical Intensive Care Unit (MICU)      16729
Medical/Surgical Intensive Care Unit (MICU/SICU)  13421
Cardiac Vascular Intensive Care Unit (CVICU)      12169
Surgical Intensive Care Unit (SICU)      11765
Trauma SICU (TSICU)      9165
Coronary Care Unit (CCU)      8746
Neuro Surgical Intensive Care Unit (Neuro SICU)      1851
Neuro Intermediate      1823
Neuro Stepdown      871
Name: first_careunit, dtype: int64
```

```
[152]: # Create admission-ICU matrix
icu_item = pd.get_dummies(icu_list['category'].apply(pd.Series).stack()).
    ↳sum(level=0)
icu_item[icu_item >= 1] = 1
icu_item = icu_item.join(icu_list['hadm_id'], how="outer")
icu_item.head()
```

```
[152]:      Cardiac Vascular Intensive Care Unit (CVICU)  Coronary Care Unit (CCU)  \
0      0      1
1      0      1
2      0      0
3      0      1
```

4		0	0
---	--	---	---

	Medical Intensive Care Unit (MICU) \
0	0
1	0
2	0
3	0
4	0

	Medical/Surgical Intensive Care Unit (MICU/SICU)	Neuro Intermediate \
0	0	0
1	0	0
2	1	0
3	0	0
4	0	0

	Neuro Stepdown	Neuro Surgical Intensive Care Unit (Neuro SICU) \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	Surgical Intensive Care Unit (SICU)	Trauma SICU (TSICU)	hadm_id
0	0	0	20000094
1	0	0	20000147
2	0	0	20000351
3	0	0	20000397
4	1	0	20000808

```
[153]: # Merge ICU data with main dataframe
final_df = admits_patients_diag.merge(icu_item, how='outer', on='hadm_id')
final_df.head()
```

```
[153]: Unnamed: 0  subject_id  hadm_id      admittime      disctime \
0          0.0  14679932.0  21038362  2139-09-26  14:16:00  2139-09-28  11:30:00
1          1.0  15585972.0  24941086  2123-10-07  23:56:00  2123-10-12  11:22:00
2          2.0  15078341.0  23272159  2122-08-28  08:48:00  2122-08-30  12:32:00
3          3.0  17301855.0  29732723  2140-06-06  14:23:00  2140-06-08  14:25:00
4          4.0  17991012.0  24298836  2181-07-10  20:28:00  2181-07-12  15:49:00
```

	deathtime	admission_type	insurance	ethnicity \
0	NaT	ELECTIVE	Other	OTHER/UNKNOWN
1	NaT	ELECTIVE	Other	WHITE
2	NaT	ELECTIVE	Other	BLACK/AFRICAN AMERICAN
3	NaT	ELECTIVE	Other	WHITE
4	NaT	ELECTIVE	Other	WHITE

```

died_at_the_hospital ...      los \
0          0.0 ...  1.884722
1          0.0 ...  4.476389
2          0.0 ...  2.155556
3          0.0 ...  2.001389
4          0.0 ...  1.806250

Cardiac Vascular Intensive Care Unit (CVICU) Coronary Care Unit (CCU) \
0          NaN          NaN
1          NaN          NaN
2          NaN          NaN
3          NaN          NaN
4          NaN          NaN

Medical Intensive Care Unit (MICU) \
0          NaN
1          NaN
2          NaN
3          NaN
4          NaN

Medical/Surgical Intensive Care Unit (MICU/SICU) Neuro Intermediate \
0          NaN          NaN
1          NaN          NaN
2          NaN          NaN
3          NaN          NaN
4          NaN          NaN

Neuro Stepdown Neuro Surgical Intensive Care Unit (Neuro SICU) \
0          NaN          NaN
1          NaN          NaN
2          NaN          NaN
3          NaN          NaN
4          NaN          NaN

Surgical Intensive Care Unit (SICU) Trauma SICU (TSICU)
0          NaN          NaN
1          NaN          NaN
2          NaN          NaN
3          NaN          NaN
4          NaN          NaN

[5 rows x 40 columns]
```

[154]:

```
# Replace NaNs with 0
final_df['Cardiac Vascular Intensive Care Unit (CVICU)'].fillna(value=0, inplace=True)
final_df['Coronary Care Unit (CCU)'].fillna(value=0, inplace=True)
final_df['Medical Intensive Care Unit (MICU)'].fillna(value=0, inplace=True)
final_df['Medical/Surgical Intensive Care Unit (MICU/SICU)'].fillna(value=0, inplace=True)
final_df['Neuro Intermediate'].fillna(value=0, inplace=True)
final_df['Neuro Stepdown'].fillna(value=0, inplace=True)
final_df['Neuro Surgical Intensive Care Unit (Neuro SICU)'].fillna(value=0, inplace=True)
final_df['Surgical Intensive Care Unit (SICU)'].fillna(value=0, inplace=True)
final_df['Trauma SICU (TSICU)'].fillna(value=0, inplace=True)
```

```
[155]: final_df.head()
```

```
[155]: Unnamed: 0  subject_id  hadm_id      admittime      disctime \
0          0.0  14679932.0  21038362  2139-09-26  14:16:00  2139-09-28  11:30:00
1          1.0  15585972.0  24941086  2123-10-07  23:56:00  2123-10-12  11:22:00
2          2.0  15078341.0  23272159  2122-08-28  08:48:00  2122-08-30  12:32:00
3          3.0  17301855.0  29732723  2140-06-06  14:23:00  2140-06-08  14:25:00
4          4.0  17991012.0  24298836  2181-07-10  20:28:00  2181-07-12  15:49:00

      deathtime admission_type insurance      ethnicity \
0          NaT      ELECTIVE      Other      OTHER/UNKNOWN
1          NaT      ELECTIVE      Other              WHITE
2          NaT      ELECTIVE      Other  BLACK/AFRICAN AMERICAN
3          NaT      ELECTIVE      Other              WHITE
4          NaT      ELECTIVE      Other              WHITE

      died_at_the_hospital  ...      los \
0                0.0  ...  1.884722
1                0.0  ...  4.476389
2                0.0  ...  2.155556
3                0.0  ...  2.001389
4                0.0  ...  1.806250

      Cardiac Vascular Intensive Care Unit (CVICU)  Coronary Care Unit (CCU) \
0                0.0                0.0
1                0.0                0.0
2                0.0                0.0
3                0.0                0.0
4                0.0                0.0

      Medical Intensive Care Unit (MICU) \
0                0.0
1                0.0
```

```

2          0.0
3          0.0
4          0.0

```

```

    Medical/Surgical Intensive Care Unit (MICU/SICU)  Neuro Intermediate  \
0          0.0          0.0
1          0.0          0.0
2          0.0          0.0
3          0.0          0.0
4          0.0          0.0

```

```

    Neuro Stepdown  Neuro Surgical Intensive Care Unit (Neuro SICU)  \
0          0.0          0.0
1          0.0          0.0
2          0.0          0.0
3          0.0          0.0
4          0.0          0.0

```

```

    Surgical Intensive Care Unit (SICU)  Trauma SICU (TSICU)
0          0.0          0.0
1          0.0          0.0
2          0.0          0.0
3          0.0          0.0
4          0.0          0.0

```

[5 rows x 40 columns]

1.4 5.Data Cleaning

```
[156]: final_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 362995 entries, 0 to 362994
Data columns (total 40 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            335257 non-null  float64
1   subject_id                            335257 non-null  float64
2   hadm_id                               362995 non-null  int64
3   admittance                             335257 non-null
datetime64[ns]
4   dischtime                             335257 non-null
datetime64[ns]
5   deathtime                             5605 non-null
datetime64[ns]
6   admission_type                         335257 non-null  object
7   insurance                             335257 non-null  object
8   ethnicity                             335257 non-null  object

```

```

9   died_at_the_hospital      335257 non-null float64
10  gender                    335257 non-null object
11  anchor_age                335257 non-null object
12  dod                       24581 non-null object
13  blood                     335257 non-null float64
14  circulatory               335257 non-null float64
15  congenital                335257 non-null float64
16  digestive                 335257 non-null float64
17  endocrine                 335257 non-null float64
18  genitourinary             335257 non-null float64
19  infectious                 335257 non-null float64
20  injury                    335257 non-null float64
21  mental                    335257 non-null float64
22  misc                      335257 non-null float64
23  muscular                  335257 non-null float64
24  neoplasms                 335257 non-null float64
25  nervous                   335257 non-null float64
26  pregnancy                 335257 non-null float64
27  prenatal                  335257 non-null float64
28  respiratory               335257 non-null float64
29  skin                      335257 non-null float64
30  los                       335257 non-null float64
31  Cardiac Vascular Intensive Care Unit (CVICU) 362995 non-null float64
32  Coronary Care Unit (CCU)   362995 non-null float64
33  Medical Intensive Care Unit (MICU)           362995 non-null float64
34  Medical/Surgical Intensive Care Unit (MICU/SICU) 362995 non-null float64
35  Neuro Intermediate         362995 non-null float64
36  Neuro Stepdown             362995 non-null float64
37  Neuro Surgical Intensive Care Unit (Neuro SICU) 362995 non-null float64
38  Surgical Intensive Care Unit (SICU)           362995 non-null float64
39  Trauma SICU (TSICU)         362995 non-null float64
dtypes: datetime64[ns](3), float64(30), int64(1), object(6)
memory usage: 113.5+ MB

```

```

[157]: # Drop unused or no longer needed columns
final_df.drop(columns=['Unnamed: 0', 'subject_id', 'hadm_id', 'admittime',
↳ 'dischtime', 'deathtime', 'dod', 'los'], inplace=True)

final_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 362995 entries, 0 to 362994
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   admission_type                        335257 non-null object
1   insurance                            335257 non-null object
2   ethnicity                            335257 non-null object

```


3	died_at_the_hospital	335257	non-null	float64
4	gender	335257	non-null	object
5	anchor_age	335257	non-null	object
6	blood	335257	non-null	float64
7	circulatory	335257	non-null	float64
8	congenital	335257	non-null	float64
9	digestive	335257	non-null	float64
10	endocrine	335257	non-null	float64
11	genitourinary	335257	non-null	float64
12	infectious	335257	non-null	float64
13	injury	335257	non-null	float64
14	mental	335257	non-null	float64
15	misc	335257	non-null	float64
16	muscular	335257	non-null	float64
17	neoplasms	335257	non-null	float64
18	nervous	335257	non-null	float64
19	pregnancy	335257	non-null	float64
20	prenatal	335257	non-null	float64
21	respiratory	335257	non-null	float64
22	skin	335257	non-null	float64
23	Cardiac Vascular Intensive Care Unit (CVICU)	362995	non-null	float64
24	Coronary Care Unit (CCU)	362995	non-null	float64
25	Medical Intensive Care Unit (MICU)	362995	non-null	float64
26	Medical/Surgical Intensive Care Unit (MICU/SICU)	362995	non-null	float64
27	Neuro Intermediate	362995	non-null	float64
28	Neuro Stepdown	362995	non-null	float64
29	Neuro Surgical Intensive Care Unit (Neuro SICU)	362995	non-null	float64
30	Surgical Intensive Care Unit (SICU)	362995	non-null	float64
31	Trauma SICU (TSICU)	362995	non-null	float64

dtypes: float64(27), object(5)
memory usage: 91.4+ MB

```
[160]: # Create dummy columns for categorical variables
prefix_cols = ['ADM', 'INS', 'ETH', 'AGE']
dummy_cols = ['admission_type', 'insurance', 'ethnicity', 'anchor_age']
df_cleaned = pd.get_dummies(final_df, prefix=prefix_cols, columns=dummy_cols)
df_cleaned.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 362995 entries, 0 to 362994
Data columns (total 44 columns):
```

#	Column	Non-Null Count	Dtype
0	died_at_the_hospital	335257 non-null	float64
1	gender	335257 non-null	object
2	blood	335257 non-null	float64
3	circulatory	335257 non-null	float64
4	congenital	335257 non-null	float64

5	digestive	335257	non-null	float64
6	endocrine	335257	non-null	float64
7	genitourinary	335257	non-null	float64
8	infectious	335257	non-null	float64
9	injury	335257	non-null	float64
10	mental	335257	non-null	float64
11	misc	335257	non-null	float64
12	muscular	335257	non-null	float64
13	neoplasms	335257	non-null	float64
14	nervous	335257	non-null	float64
15	pregnancy	335257	non-null	float64
16	prenatal	335257	non-null	float64
17	respiratory	335257	non-null	float64
18	skin	335257	non-null	float64
19	Cardiac Vascular Intensive Care Unit (CVICU)	362995	non-null	float64
20	Coronary Care Unit (CCU)	362995	non-null	float64
21	Medical Intensive Care Unit (MICU)	362995	non-null	float64
22	Medical/Surgical Intensive Care Unit (MICU/SICU)	362995	non-null	float64
23	Neuro Intermediate	362995	non-null	float64
24	Neuro Stepdown	362995	non-null	float64
25	Neuro Surgical Intensive Care Unit (Neuro SICU)	362995	non-null	float64
26	Surgical Intensive Care Unit (SICU)	362995	non-null	float64
27	Trauma SICU (TSICU)	362995	non-null	float64
28	ADM_ELECTIVE	362995	non-null	uint8
29	ADM_EMERGENCY	362995	non-null	uint8
30	ADM_OBSERVATION	362995	non-null	uint8
31	ADM_SURGICAL SAME DAY ADMISSION	362995	non-null	uint8
32	INS_Medicaid	362995	non-null	uint8
33	INS_Medicare	362995	non-null	uint8
34	INS_Other	362995	non-null	uint8
35	ETH_ASIAN	362995	non-null	uint8
36	ETH_BLACK/AFRICAN AMERICAN	362995	non-null	uint8
37	ETH_HISPANIC/LATINO	362995	non-null	uint8
38	ETH_OTHER/UNKNOWN	362995	non-null	uint8
39	ETH_WHITE	362995	non-null	uint8
40	AGE_MIDDLE_ADULT	362995	non-null	uint8
41	AGE_NEWBORN	362995	non-null	uint8
42	AGE_SENIOR	362995	non-null	uint8
43	AGE_YOUNG_ADULT	362995	non-null	uint8

dtypes: float64(27), object(1), uint8(16)

memory usage: 85.9+ MB

```
[172]: # Drop rows that contain NaN values
df_cleaned.dropna(axis=0, inplace=True)
```

```
[175]: # Check for any remaining NaNs
df_cleaned.isnull().values.sum()
```

[175]: 0

```
[173]: df_cleaned.isnull().sum()
```

```
[173]: died_at_the_hospital      0
gender                        0
blood                        0
circulatory                  0
congenital                   0
digestive                    0
endocrine                    0
genitourinary                0
infectious                   0
injury                       0
mental                       0
misc                         0
muscular                     0
neoplasms                    0
nervous                      0
pregnancy                    0
prenatal                     0
respiratory                  0
skin                         0
Cardiac Vascular Intensive Care Unit (CVICU) 0
Coronary Care Unit (CCU)      0
Medical Intensive Care Unit (MICU) 0
Medical/Surgical Intensive Care Unit (MICU/SICU) 0
Neuro Intermediate           0
Neuro Stepdown               0
Neuro Surgical Intensive Care Unit (Neuro SICU) 0
Surgical Intensive Care Unit (SICU) 0
Trauma SICU (TSICU)          0
ADM_ELECTIVE                 0
ADM_EMERGENCY                0
ADM_OBSERVATION              0
ADM_SURGICAL SAME DAY ADMISSION 0
INS_Medicaid                0
INS_Medicare                 0
INS_Other                    0
ETH_ASIAN                    0
ETH_BLACK/AFRICAN AMERICAN  0
ETH_HISPANIC/LATINO          0
ETH_OTHER/UNKNOWN            0
ETH_WHITE                    0
AGE_MIDDLE_ADULT             0
AGE_NEWBORN                  0
AGE_SENIOR                   0
```

```
AGE_YOUNG_ADULT
dtype: int64
```

```
[176]: # Convert gender into numeric boolean attribute
df_cleaned['gender'].replace({'M': 0, 'F': 1}, inplace=True)
df_cleaned.head()
```

```
[176]:
```

	died_at_the_hospital	gender	blood	circulatory	congenital	digestive	\
0	0.0	1	0.0	0.0	1.0	0.0	
1	0.0	1	0.0	0.0	0.0	0.0	
2	0.0	0	0.0	0.0	0.0	0.0	
3	0.0	1	0.0	0.0	0.0	0.0	
4	0.0	0	0.0	0.0	0.0	0.0	

	endocrine	genitourinary	infectious	injury	...	INS_Other	ETH_ASIAN	\
0	0.0	0.0	0.0	2.0	...	1	0	
1	0.0	0.0	0.0	2.0	...	1	0	
2	0.0	0.0	0.0	3.0	...	1	0	
3	0.0	0.0	0.0	2.0	...	1	0	
4	0.0	0.0	0.0	2.0	...	1	0	

	ETH_BLACK/AFRICAN AMERICAN	ETH_HISPANIC/LATINO	ETH_OTHER/UNKNOWN	\
0	0	0	1	
1	0	0	0	
2	1	0	0	
3	0	0	0	
4	0	0	0	

	ETH_WHITE	AGE_MIDDLE_ADULT	AGE_NEWBORN	AGE_SENIOR	AGE_YOUNG_ADULT
0	0	0	1	0	0
1	1	0	1	0	0
2	0	0	1	0	0
3	1	0	1	0	0
4	1	0	1	0	0

[5 rows x 44 columns]

```
[177]: df_cleaned = df_cleaned.astype(int)
```

```
[178]: # Check for any remaining NaNs
df_cleaned.isnull().values.sum()
```

```
[178]: 0
```

1.5 6. Prediction Model

We use a **Supervised Learning ML model**. First of all what is it? Supervised learning is defined by its use of labeled datasets to train algorithms that to classify data or predict outcomes accurately.

It uses a training set to teach models to yield the desired output. This training dataset includes inputs and correct outputs, which allow the model to learn over time. The algorithm measures its accuracy through the loss function, adjusting until the error has been sufficiently minimized.

Why do we choose it? Because in our case we have the correct output for each dataset entry: “Died_at_the_Hospital” (Yes or No) and we want to create a model that predicts this output for new entries, in other words that it “generalize well”.

We will implement the supervised learning prediction model using the **Scikit-Learn** machine learning library.

To implement the prediction model, our dataset is splitted into training and test sets at an 80:20 ratio using the scikit-learn *train_test_split* function.

Why split in training and test set? Because to detect a machine learning model behavior, we need to use observations that aren’t used in the training process. Otherwise, the evaluation of the model would be biased as a matter of fact when we build a predictive model, we want the model to work well on data that the model has never seen, so that’s the reason why we use a training set to train the model and a test set to evaluate the model accuracy.

Searching on the Internet for the best train-test ratio, the first answer is 80:20. This means we use 80% of the observations for training and the rest for testing. This approach is taken in this case. zability)

```
[179]: # Target Variable (died_at_the_hospital)
HOSP_MORT = df_cleaned['died_at_the_hospital'].values
# Prediction Features
features = df_cleaned.drop(columns=['died_at_the_hospital'])

[183]: from sklearn.model_selection import train_test_split

# Split into training set 80% and test set 20%
X_train, X_test, y_train, y_test = train_test_split(features,
                                                    HOSP_MORT,
                                                    test_size = .20,
                                                    random_state = 0)

# Show the results of the split
print("Training set has {} samples.".format(X_train.shape[0]))
print("Testing set has {} samples.".format(X_test.shape[0]))
```

Training set has 268205 samples.

Testing set has 67052 samples.

Using the training set, we'll four five different classification models (from the scikit-learn library) using default settings.

```
[186]: from sklearn.metrics import r2_score
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
```

```

from sklearn.ensemble import GradientBoostingRegressor

# Regression models used from scikit-learn for comparison
models = [GradientBoostingRegressor(random_state = 0),
          LinearRegression(),
          KNeighborsRegressor(),
          RandomForestRegressor(random_state = 0)]

results = {}

for model in models:
    # Instantiate and fit Regressor Model
    reg_model = model
    reg_model.fit(X_train, y_train)

    # Make predictions with model
    y_test_preds = reg_model.predict(X_test)

    # Grab model name and store results associated with model
    name = str(model).split("(")[0]

    results[name] = r2_score(y_test, y_test_preds)
    print('{} done.'.format(name))

```

GradientBoostingRegressor done.

LinearRegression done.

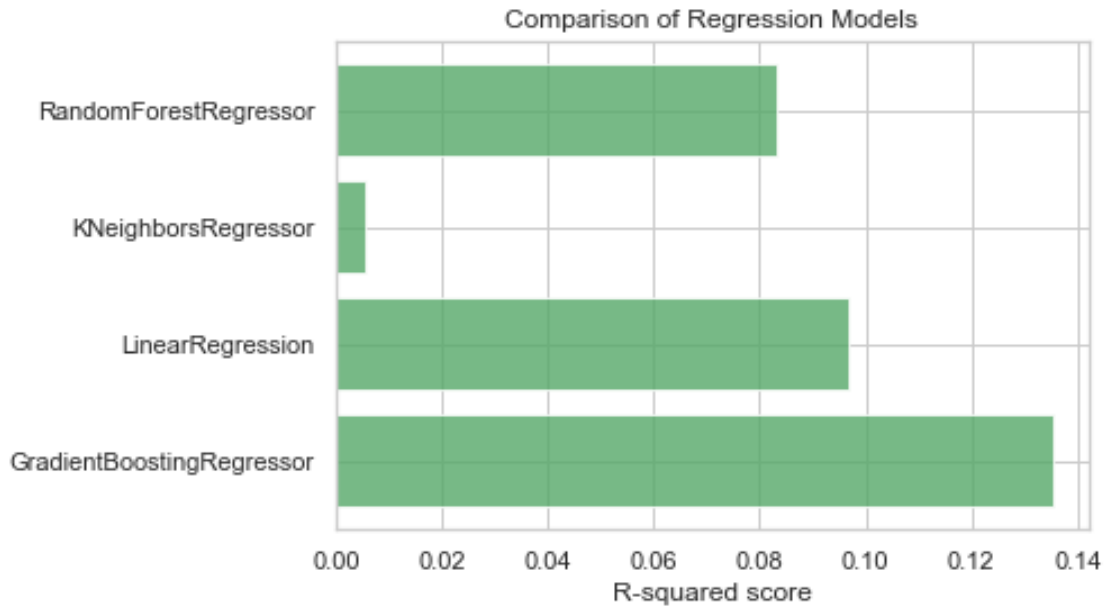
KNeighborsRegressor done.

RandomForestRegressor done.

```

[203]: # R2 score results
fig, ax = plt.subplots()
ind = range(len(results))
ax.barh(ind, list(results.values()), align='center',
        color = '#55a868', alpha=0.8)
ax.set_yticks(ind)
ax.set_yticklabels(results.keys())
ax.set_xlabel('R-squared score')
ax.tick_params(left=False, top=False, right=False)
ax.set_title('Comparison of Regression Models')
fig.savefig('images/compare_models_dinh_mimic4.png', bbox_inches = 'tight')

```



```
[192]: # GradientBoostingRegressor will be used as the In-Hospital mortality
        ↳ prediction model
        # GradientBoostingRegressor will be used as the LOS prediction model
        reg_model = GradientBoostingRegressor(random_state=0)
        reg_model.fit(X_train, y_train)
        y_test_preds = reg_model.predict(X_test)
        r2_not_refined = r2_score(y_test, y_test_preds)
        print("R2 score is: {:.2f}".format(r2_not_refined))
```

R2 score is: 0.135291

The GradientBoostingRegressor has the best R2 score of ~13% so we focus on refining this particular model.

1.6 7. Parameter Tuning

To refine the GradientBoostingRegressor model, **GridSearchCV** function from scikit-learn is used to test out various permutations of parameters such as *n_estimators*, *max_depth*, and *loss*. It helps to loop through predefined hyperparameters and fit your estimator (model) on your training set. So, in the end, we could select the best parameters from the listed hyperparameters.

```
[197]: from sklearn.model_selection import GridSearchCV

        # Split into train 80% and test 20%
        X_train, X_test, y_train, y_test = train_test_split(features,
                                                            HOSP_MORT,
                                                            test_size = .20,
                                                            random_state = 42)
```

```

# Set the parameters by cross-validation
#tuned_parameters = [{'n_estimators': [100, 200, 300],
#                      'max_depth' : [2, 3, 4],
#                      'loss': ['ls', 'lad', 'huber']}]]
tuned_parameters = [{'n_estimators': [200, 300],
                      'max_depth' : [3, 4],
                      'loss': ['ls', 'lad']}]]

# create and fit a ridge regression model, testing each alpha
reg_model = GradientBoostingRegressor()
grid = GridSearchCV(reg_model, tuned_parameters, cv=3, verbose = 1, n_jobs = -1)
grid.fit(X_train, y_train)
reg_model_optimized = grid.best_estimator_

# summarize the results of the grid search
print(grid.best_score_)
print(grid.best_estimator_)

```

Fitting 3 folds for each of 8 candidates, totalling 24 fits

0.14951616540043464

GradientBoostingRegressor(max_depth=4, n_estimators=200)

Tuned Paramters - *n_estimators*: The number of boosting stages to perform. - *max_depth*: maximum depth of the individual regression estimators. The maximum depth limits the number of nodes in the tree. - *loss*: loss function to be optimized. 'ls' refers to least squares regression. 'lad' (least absolute deviation) is a highly robust loss function solely based on order information of the input variables. 'huber' is a combination of the two.

The best estimator result from GridSearchCV was n_estimators=200, max_depth=4, loss = ls.

```

[198]: y_test_preds = reg_model_optimized.predict(X_test)
r2_optimized = r2_score(y_test, y_test_preds)
print("Optimized R2 score is: {:.2f}".format(r2_optimized))

```

Optimized R2 score is: 0.143552

```

[199]: print('Parameter tuning improved R2 score by {:.4f}'.
        ↳format(r2_optimized-r2_not_refined))

```

Parameter tuning improved R2 score by 0.0083

1.7 8. Model evaluation and result Discussion

We could look at what features were most important in predicting in-hospital mortality when using the gradient boosting regression model.

```

[200]: feature_imp = pd.DataFrame(reg_model_optimized.feature_importances_,
                                index = X_train.columns,

```



```

                                columns=['importance'])).
↪sort_values('importance', ascending=False)

feature_imp.head(20)

```

```


[200]:
respiratory            0.212500
Medical Intensive Care Unit (MICU)    0.086784
Surgical Intensive Care Unit (SICU)   0.076962
Medical/Surgical Intensive Care Unit (MICU/SICU) 0.076263
neoplasms              0.068504
infectious             0.053728
genitourinary          0.049491
Coronary Care Unit (CCU)             0.042526
circulatory            0.040123
Trauma SICU (TSICU)         0.036257
misc                    0.033244
digestive              0.028575
blood                  0.024942
injury                 0.022731
ETH_OTHER/UNKNOWN        0.021645
nervous                0.019831
AGE_SENIOR              0.019813
ADM_EMERGENCY           0.015731
mental                  0.011016
endocrine               0.008235

```

```

[202]: #Let's plot the top-10 feature importance
feature_imp.index[0:10].tolist()
# Plot feature importance
fig, ax = plt.subplots(figsize=(7, 5))
ind = range(0,10)
ax.barh(ind, feature_imp['importance'].values[0:10],
        align='center', color='#c44e52', alpha=0.9)
ax.set_yticks(ind)
ax.set_yticklabels(feature_imp.index[0:10].tolist())
ax.tick_params(left=False, top=False, right=False)
ax.set_title("Top 10 features for predicting LOS")
ax.set_xlabel('Feature Importance Coefficient \n(GradientBoostingRegressor)')
plt.gca().invert_yaxis()
fig.savefig('images/feature_importance_dinh_mimic4.png', bbox_inches = 'tight')

```



- We could say that, first of all, one of the results is that the *ICD-9 diagnoses categories* and the admission to various type of ICU, are by far the most important features between the features analyzed.
- We could notice how the diagnose belongig to **respiratory** category is the most important feature in determining **in-hospital-mortality** followe by the admission of the patient to MICU, SICU or both.

Compute now other metric used for validation RMSE.

```
[ ]: for i in range(y_test_preds.shape[0]):
    ml_model = abs(y_test_preds[i] - y_test[i])

    ml_model_rms = (y_test_preds[i] - y_test[i])**2

print("Prediction Model RMS {}".format((ml_model_rms**0.5)/y_test_preds.
    ↳shape[0]))
```

1.8 Conclusions for In-Hospital Mortality

The development of methods for prediction of mortality rates in populations has been motivated primarily by the need to compare the efficacy of medications, care guidelines, surgery, and other interventions when, as is common, it is necessary to control for differences in diagnoses, age, and other factors.

The prediction model achieved should be a “friend” of the doctor and not a substitute. In this case the model could help the doctor to adopt certain behaviors or to carry out different interventions depending on the expected probability of mortality in hospital.

For example, for a patient who has a high probability of death during his hospital stay, the doctor may consider adopting some surgical interventions right away without waiting for further tests and

therefore reduce the expected probability of death.

Future Developments in this area could be analyzing other features of mimic dataset or also other datasets, for example the notes of doctors using Natural Language Processing and trying to have insights.